



Levina--Bickel MLE: raw versus locally detrended (W=100)

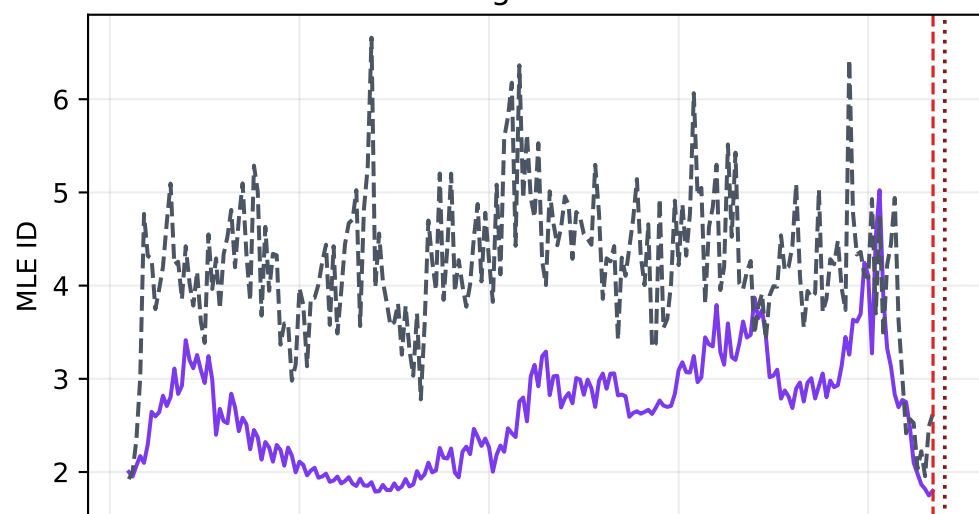
train loss



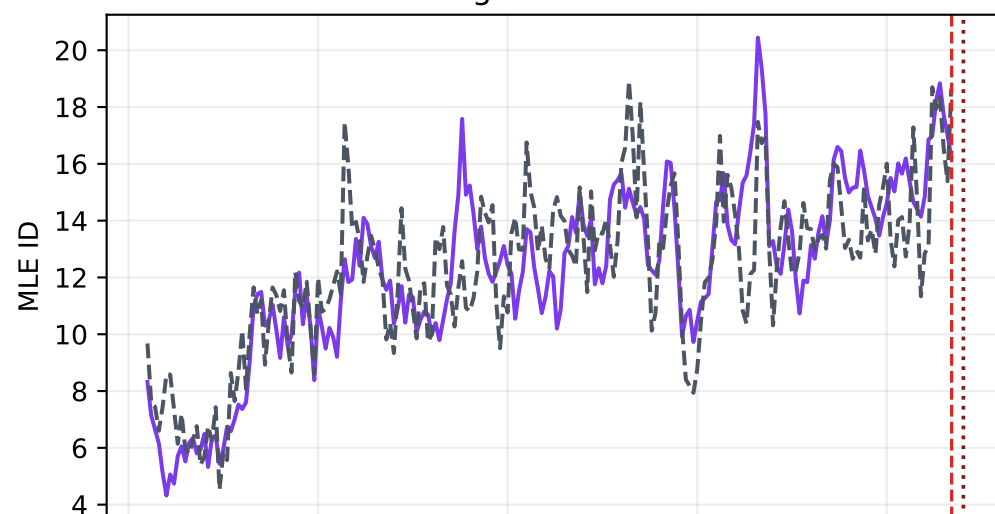
val loss



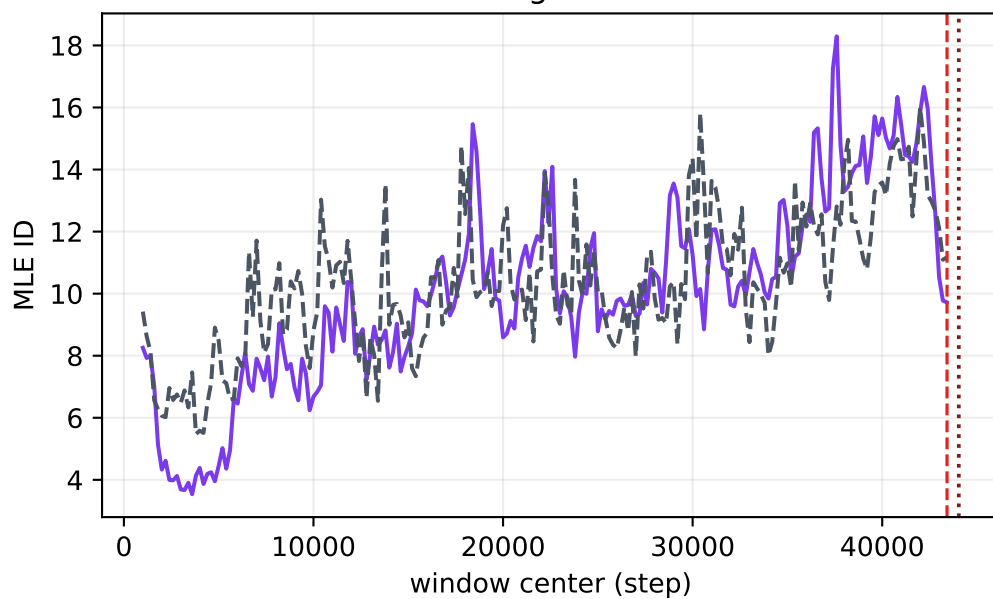
weight norm



grad norm



embed grad norm



grad cosine

