
Counting the Active Dimension of a Training Run: What a Single Scalar Log Can and Cannot Measure

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Over a stretch of training, the parameter trajectory of a neural network stays on a
2 set far smaller than the parameter space. We call the dimension of that set the active
3 dimension of the stretch. Measuring it requires storing the trajectory, whereas a
4 run writes scalar logs at no cost. We ask how much of it can be recovered from one
5 such scalar. We reconstruct a state space from the scalar by delay coordinates and
6 estimate the intrinsic dimension of that reconstruction by maximum likelihood on
7 nearest-neighbour distances. The procedure returns a plausible number even when
8 the trajectory never revisits any region of the parameter space, so most of the work
9 is in establishing when its output may be believed. On six constructed systems of
10 known active dimension it is accurate to about one component, up to a ceiling near
11 eight. Statistics computed from the log identify the records on which the value may
12 not be read as a dimension at all. They reject both standard grokking settings, so we
13 store the trajectory instead. Its effective rank collapses near the generalisation step
14 in four regularised runs and then re-expands, the simplification that mechanistic
15 accounts of grokking predict; a run that has merely come to rest falls as far: the
16 sharpness and the reversal mark the transition, not the depth. Re-run over a window
17 as short as the transition, the estimate on an ordinary log falls at the same step in
18 those four runs and in neither of the other two, by more than randomised series that
19 preserve the shape of the log.

20 1 Introduction

21 A network is trained in a parameter space of dimension P , but over a stretch of training the parameters
22 usually move within a far smaller set. Two open questions turn on that set. Is an apparent simplification
23 of training a real reduction in the number of degrees of freedom in use, or an artefact of a shrinking
24 scale? Can a phase transition such as delayed generalisation be detected without instrumenting the
25 run? Both are argued with whichever count of degrees of freedom is least expensive to compute. The
26 counts in use disagree: the *available* dimension of the subspace the parameters are confined to [Li
27 et al., 2018], the *functional* dimension of the map from that subspace to the outputs, and the *effective*
28 *rank* of the trajectory covariance.

29 Our two questions concern a fourth quantity: the number of independent components of the set the
30 parameters occupy. We call it the *active dimension* of that stretch (equation (1)). If r independent
31 quasiperiodic phases drive the motion, the path lies on an r -dimensional torus and the active dimension
32 is r . Section 3.1 defines all four. Section 5.2 exhibits a change that moves the active dimension while
33 the other three stay where they are.

34 Measuring the active dimension means storing the parameter trajectory, which an ordinary run does
35 not do. A run does write scalar logs at no extra cost: a loss, a gradient norm, a parameter norm. *How*
36 *much of the active dimension can be recovered from one such scalar?* Nonlinear time series analysis
37 supplies a standard answer: reconstruct a state space from the scalar by delay coordinates [Takens,

1981, Sauer et al., 1991], then estimate the intrinsic dimension of the reconstruction by maximum likelihood on nearest-neighbour distances [Levina and Bickel, 2004, MacKay and Ghahramani, 2005]. Appendix A states these results.

The procedure returns a number on any record. Power-law noise yields a finite and reproducible value fixed by its spectral exponent [Osborne and Provenzale, 1989, Theiler, 1991]; a decaying transient yields a large value. Establishing which case a training log falls in is the substance of the problem. Most of the work below is spent on it, beginning with the validation protocol of section 4.

Contributions.

1. We evaluate the estimate against a constructed ground truth on six systems of increasing realism, ending with a linear head trained inside a k -dimensional subspace on top of a frozen image network. On withheld ranks and seeds it recovers the active dimension to within about one component, up to a ceiling near eight (section 5).
2. We establish the conditions the value depends on. The identifiability ratio and the trend-crossing count identify the failure regimes from the series alone (section 6).
3. We apply the procedure to grokking (section 7). Both published settings fall outside the regime the estimate needs, so we store the trajectory instead. Its effective rank collapses at generalisation and re-expands. Over a window as short as that transition the estimate on a log declines with it in the four runs that generalise and in neither of the other two (section 7.3).

2 Related work

Reconstruction and intrinsic dimension. Delay reconstruction follows from Takens’ theorem and its extensions (Appendix A), is used on empirical series [Sugihara and May, 1990, Sugihara et al., 2012], and takes its lag and dimension from mutual information, false nearest neighbours or singular spectra [Fraser and Swinney, 1986, Kennel et al., 1992, Cao, 1997, Broomhead and King, 1986]. We use pooled maximum likelihood [Levina and Bickel, 2004, MacKay and Ghahramani, 2005]; the alternatives [Grassberger and Procaccia, 1983, Facco et al., 2017, Camastra and Staiano, 2016] are bounded alike by a finite record [Eckmann and Ruelle, 1992]. Oversampling [Theiler, 1986] and power-law noise [Osborne and Provenzale, 1989] produce spurious dimensions; surrogate tests [Theiler et al., 1992, Schreiber and Schmitz, 1996] detect nonlinear structure but not how much is present and miss a decaying transient entirely.

Dimension in optimisation. Random-subspace training [Li et al., 2018, Aghajanyan et al., 2021] reports the smallest k that still solves a task, a property of the task and not of a window; we adopt the same parameterisation but measure a different quantity within it. Closest to the active dimension is the finding that descent occurs largely in a small subspace spanned by the top Hessian eigenvectors [Gur-Ari et al., 2018, Sagun et al., 2017]. Ansuini et al. [2019] and Pope et al. [2021] estimate the intrinsic dimension of a data set or of a layer’s representations, which are properties of a point cloud, whereas the active dimension is a property of a path. The participation ratio we adopt [Gao et al., 2017] is not the spectral-entropy quantity of the same name [Roy and Vetterli, 2007]; it has been measured on a grokking trajectory once, in passing [Notsawo Jr. et al., 2023]. All of these quantities need weights, gradients or activations, whereas the nearest precedent for a scalar computed during training and acted on is the gradient noise scale [McCandlish et al., 2018]. Mini-batch descent near a minimum is modelled as a stochastic process [Mandt et al., 2017]; Appendix K asks whether undriven descent at the edge of stability [Cohen et al., 2021] returns near states it has already visited.

Grokking. Delayed generalisation [Power et al., 2022] is studied in two settings and we use both. The transformer setting is regularised and, in our version, stochastic: the configuration of Nanda et al. [2023], also used by Liu et al. [2023], on modular addition and on the S_5 task of Power et al. [2022]. The perceptron setting is the unregularised full-batch gradient descent of Gromov [2023], on the modular polynomials and representability criterion of Doshi et al. [2024], which supply label-matched pairs (Appendix J). Explanations are framed in weight or activation space [Nanda et al., 2023, Varma et al., 2023, Liu et al., 2022, Barak et al., 2022] or in a lazy-to-rich transition [Kumar et al., 2024]. Closest to us, Notsawo Jr. et al. [2023] predict grokking from early loss oscillations [Thilak et al., 2022] and Merrill et al. [2023] track competing subnetworks. We differ from all of these in estimating a dynamical quantity whose ground truth is measured independently.

90 3 Problem statement

91 3.1 The active dimension

92 A *window* $\mathcal{W}$ is a block of L consecutive optimiser steps. Every quantity below is defined on a single
 93 window. A run is described by a sequence of them. Let $\theta_t \in \mathbb{R}^P$ be the parameter vector at step t ,
 94 confined by construction or by the dynamics to an affine subspace $\theta = \theta_0 + V^\top c$ with $V \in \mathbb{R}^{k \times P}$
 95 orthonormal, so that $c_t \in \mathbb{R}^k$ holds the free coordinates.

96 We call k the *available* dimension. The *functional* dimension is the rank of $\partial f(c)/\partial c$, with f the
 97 outputs on a fixed probe set; it bounds how many of the k directions can change the computed
 98 function. The *active* dimension is the subject of this paper: the number of independent directions
 99 along which c_t moves within the window. Let $\mathcal{M}_{\mathcal{W}} = \{c_t : t \in \mathcal{W}\}$ be the set of points it occupies
 100 there. The active dimension is the box-counting dimension of $\mathcal{M}_{\mathcal{W}}$ at resolution ε ,

$$d_{\text{act}}(\mathcal{W}) = \dim_{\varepsilon} \mathcal{M}_{\mathcal{W}}. \quad (1)$$

101 Here ε *stays finite*, whereas a box-counting dimension is ordinarily defined as a limit $\varepsilon \rightarrow 0$. It must
 102 stay finite because the estimator reaches $\mathcal{M}_{\mathcal{W}}$ only through distances between nearby samples. *The*
 103 *window must also be long enough*: any path resembles a curve over a short stretch, giving $d_{\text{act}} = 1$
 104 unless $\mathcal{W}$ outlasts the return time of the trajectory.

105 The systems of section 5 set $c_t = F(\varphi_1(t), \dots, \varphi_r(t))$, with $\varphi_1, \dots, \varphi_r$ independent quasiperiodic
 106 phases. Because the phases are rationally independent, the path lies on an r -dimensional torus and
 107 passes arbitrarily close to every point of it, so $d_{\text{act}} = r$ at every ε below the amplitude of the weakest
 108 phase. Above that scale the weakest phases are invisible and d_{act} declines with the resolution, as the
 109 definition requires (section 6.4).

110 We also measure the participation ratio [Gao et al., 2017], which for a non-negative spectrum
 111 $\lambda_1, \dots, \lambda_n$ is

$$\text{PR}(\lambda) = \frac{(\sum_i \lambda_i)^2}{\sum_i \lambda_i^2}. \quad (2)$$

112 It equals m exactly when m of the λ_i are equal and the rest vanish. Applied to the eigenvalues of a
 113 covariance over the window it gives the *effective rank*, in two forms named by the superscript:

$$\text{PR}^{\text{pos}} = \text{PR}(\text{Cov} c_t), \quad \text{PR}^{\text{det}} = \text{PR}(\text{Cov} \tilde{c}_t), \quad (3)$$

114 both over $t \in \mathcal{W}$, with $\tilde{c}_t$ the residual of c_t about its per-window least-squares line. A drifting
 115 trajectory needs the detrended form, since a straight drift dominates the spectrum and holds PR^{pos}
 116 near unity however the trajectory moves about that line. Written without a superscript, *the effective*
 117 *rank* means PR^{pos} .

118 The effective rank is not the active dimension. It equals r when r directions carry comparable
 119 variance and falls below r when they do not, while d_{act} stays at r (section 6.4). We use it to check
 120 that each system of section 5 really does excite r independent phases (table 11). It is also the quantity
 121 section 7.2 measures on a stored grokking trajectory, in the setting where a log yields no dimension
 122 at all.

123 3.2 Observation model and estimator

124 We observe a scalar series $x_t = \phi(s_t)$, in which s_t is the full optimiser state and ϕ is fixed before the
 125 run. From it we reconstruct $y_t = (x_t, \dots, x_{t-(E-1)\tau}) \in \mathbb{R}^E$. For a generic pair of dynamics and
 126 observation on a compact invariant set of box-counting dimension d , this map is an embedding once
 127 $E > 2d$ (theorems 1 and 2). Genericity is not verifiable for a ϕ fixed in advance, so every result below
 128 is reported over a family of observers. We estimate the intrinsic dimension of $\{y_t\}$ by maximum
 129 likelihood on nearest-neighbour distances (theorem 4). Let $r_1^{(i)} \leq \dots \leq r_m^{(i)}$ be the distances from
 130 point i to its m nearest neighbours, where a neighbour counts only if it is more than W_T steps away
 131 from i in time. Discarding the near-in-time neighbours in this way is the *Theiler exclusion* [Theiler,
 132 1986, Kantz and Schreiber, 2004]: without it, the nearest points in $\mathbb{R}^E$ are the samples just before
 133 and after i , so the statistic measures the sampling rate instead of the geometry. Pooling the likelihood

over the N points before inverting, as MacKay and Ghahramani [2005] recommend, gives

$$\hat{d}_{\text{MG}}^{-1} = \frac{1}{N(m-1)} \sum_{i=1}^N \sum_{j=1}^{m-1} \log \frac{r_m^{(i)}}{r_j^{(i)}}, \quad (4)$$

while averaging the per-point estimates instead, as Levina and Bickel [2004] propose, gives the variant we write LB. We write *the estimate* for the value MG returns and *the estimator* for the procedure; table 8 reports both and Appendix B.1 gives the algorithm as implemented. The *roughness* $\text{std}(\Delta x)/\text{std}(x)$, the *linear participation ratio* PR_{delay} of equation (2) on $\text{Cov}\{\mathbf{y}_t\}$ and the *spectral participation ratio* on the power spectrum of x accompany every estimate. Each can grow with r for reasons that involve no geometry. If one of them recovers the truth as well as the estimator does, the accuracy of the estimator is no evidence that a dimension has been measured.

3.3 Dynamical regimes

Two conditions must hold before an estimate computed from a log may be read as an active dimension. The embedding theorems (theorems 1 and 2) require a deterministic map on a compact invariant set; that licenses the delay reconstruction as a change of coordinates. The neighbour statistic of equation (4) requires more: the record must cover that set densely enough that a point's nearest neighbours in $\mathbb{R}^E$ are returns to its neighbourhood rather than the samples immediately before and after it. The Theiler exclusion enforces this and section 3.4 tests it. The estimator itself would accept points drawn from independent runs; recurrence is needed because our entire sample is one trajectory. An optimiser near a solution may be in any of three regimes. Only the first satisfies both conditions.

- **Deterministic and recurrent.** The path returns close to states already passed through and lies on an r -dimensional torus, so the occupied set has r independent components and $d_{\text{act}} = r$. Here both conditions hold and the estimator recovers r . This is the only regime in which it does.
- **Deterministic and transient.** A converging run decays along a curve it never revisits, so the occupied set is that curve and $d_{\text{act}} = 1$. The estimand exists, but the protocol misses it. The record contains no returns: in place of 1 the estimator reports about 29 at every r , a value produced by the Theiler exclusion itself (section 6.1). The effective rank does stay near unity here (table 11), but only because the curve is nearly straight.
- **Stochastically driven.** Mini-batch gradient noise leaves no low-dimensional invariant set for the path to fill, so there is no active dimension to estimate; the value returned is set by E_{max} and by the record length. The embedding theorems for forced systems (theorem 3) do cover a drive of this kind, but one realisation of it at a time, whereas a single run crosses realisations.

3.4 Diagnostics

The series alone is enough to tell a transient from a stochastic drive. The **identifiability ratio** $\rho_{\text{ident}} = \hat{d}_{\text{MG}}(2E_{\text{max}})/\hat{d}_{\text{MG}}(E_{\text{max}})$ compares the estimate at a doubled embedding dimension with the estimate at the working one. It stays near unity where a dimension can be measured. Under mini-batch noise the estimate follows the embedding dimension rather than the data; the ratio then grows. A transient escapes it: there the estimate is stable in the embedding dimension for reasons unconnected with any dimension. A transient is caught by the **trend-crossing count** instead, the number of sign changes of the residual of the window about its own least-squares line, which is near zero on a monotone decay. The count tests only non-monotonicity, so it cannot replace the ratio; we report both with every estimate. Section 6.4 validates both.

4 Validation protocol

Four requirements were fixed before any result was computed. Table 4 records which of them each system fails.

Constructed truth: the ground truth is measured rather than assumed. Every system is built with r independent quasiperiodic phases, so that $d_{\text{act}} = r$ by construction. The effective rank of the recorded trajectory must then confirm that all r are comparably excited.

179 *Frozen configuration:* the estimator is chosen once, by a search over its own settings alone. Searching
 180 the settings of the system at the same time would select the data together with the method: the result
 181 would then report how favourable a system the search had found.

182 *Withheld ranks:* ranks are held out as well as seeds, because the geometry of the excitation depends
 183 on r and every seed at a given rank shares that geometry.

184 *Zero learning rate:* every system is re-run with the learning rate set to zero, so that the parameters
 185 never move and any observer of the optimiser must be constant. A log that still varies and still returns
 186 its earlier estimate was never a function of the optimiser. The frozen nonlinear decoder and the
 187 perceptron in a k -subspace fail it.

188 5 Evaluation on systems with a known active dimension

189 We evaluate the estimator on six systems of increasing realism: an oscillating diagonal matrix,
 190 online linear regression, logistic regression, a frozen nonlinear decoder, a perceptron confined to a
 191 k -dimensional subspace, and a linear head on a network trained on image data. In each, the path lies
 192 on an r -dimensional torus by construction, giving $d_{\text{act}} = r$; the measured effective rank confirms
 193 that all r phases are excited; the estimator receives one scalar series and nothing else. Table 4 collects
 194 the outcome; the simplest and the most realistic are described below.

195 5.1 An oscillating diagonal matrix

196 The first system has no optimiser and no feedback. In $\mathbf{W}(t) = \text{diag}(b_i + \delta_i(t))$, the offsets b_i are
 197 fixed and r of the δ_i oscillate at rationally independent frequencies. The estimate tracks r across the
 198 whole withheld range to within a third of a component (table 4).

199 5.2 A constrained head on a network trained on image data

200 A tanh network is trained on the scikit-learn handwritten digits and then frozen. A linear head
 201 confined to a k -dimensional affine subspace is trained on top of it under softmax cross-entropy, so that
 202 the curvature is data-dependent and all the dynamics occur in $\mathbf{c}$. The excitation is made r -dimensional
 203 by modulating the loss weights of r of twelve fixed data groups and orthogonalising the resulting
 204 gradient directions (Appendix E). We drive the head quasiperiodically, with mini-batch noise, and by
 205 plain gradient descent, covering the recurrent, stochastic and transient regimes of section 3.3.

206 Twelve scalar observers are recorded (Appendix B.2). The estimator was frozen on four of them, on
 207 seeds and ranks reserved for that purpose. Every result below is therefore out of sample in rank, in
 208 seed, and for most observers in the observer too. In the deterministic recurrent regime the estimate
 209 recovers the active dimension to within about one component, though the error differs by a factor
 210 of five between the best observer and the worst (table 6). Every observer orders the withheld ranks
 211 correctly (table 4). For none of the ten scored observers does the roughness of the log order those
 212 ranks, so the agreement cannot be explained by the logs merely becoming rougher as r grows. At
 213 $k = 20$ the same system behaves alike at a lower ceiling (table 7). We mark that result exploratory,
 214 because it was scored with a Theiler exclusion smaller than the protocol requires (Appendix I.1).

215 **The ceiling.** The estimator saturates above about eight components. The embedding dimension sets
 216 the level it returns at high r , while the largest r it still tracks follows neither the embedding dimension
 217 nor the record length and never exceeds about eleven (Appendix L). One caveat is substantial. The
 218 participation ratio of the delay covariance, a purely linear statistic, saturates along the same curve, so
 219 the limit may be the number of components the delay window resolves linearly, with no geometry
 220 involved.

221 **The zero-learning-rate test.** The zero-learning-rate check requires that an observer of a frozen
 222 optimiser be constant. Most published analyses of dimension on training logs do not run it. It
 223 invalidated two of our systems, which reproduce their reported results with the parameters frozen.
 224 It also rejects the instantaneous mini-batch loss of section 5.2, otherwise among the more accurate
 225 observers.




Figure 1: Recovery of the active dimension in each regime. One system, the linear head of section 5.2, driven five ways: a fast and a slow quasiperiodic drive (*recurrent, fast* and *recurrent, slow*), the fast drive with gradient noise added (*recurrent + noise*), mini-batch noise alone (*stochastic*), and plain gradient descent, which decays without recurring (*transient*). Colour is the regime of section 3.3. (a) Estimate against truth: marker, median over seeds and observers; band, their interquartile range; dashed diagonal, exact recovery; dotted line, $E_{\max}$. (b) The identifiability ratio on the same five, one point per observer, at $r = 2$ (open) and $r = 6$ (solid).

6 Conditions of validity

6.1 Recovery in each regime

Figure 1 evaluates the estimator in each of the three regimes of section 3.3. Its two panels ask the same question with and without the ground truth. Panel (a) shows which cases the estimator recovers correctly. Panel (b) gives the conclusion a user with no ground truth to check against would have drawn about the same cases.

In (a) the fast quasiperiodic drive follows the diagonal over the whole withheld range. Under mini-batch noise the estimate is flat in r and well above the truth, responding to $E_{\max}$ and to the length of the record instead of to the data. Weak noise on an otherwise recurrent trajectory leaves the ordering intact and destroys the absolute value. The slow drive is the informative case, because it changes behaviour partway along: it tracks the truth to about $r = 4$, then saturates.

Panel (b) gives the identifiability ratio for the same cases at $r = 2$ and $r = 6$, on either side of that change. The regime labels alone would have predicted the extreme cases, since the fast drive passes at both ranks and the transient is furthest from passing at both. They would not have predicted the slow drive, which passes at $r = 2$, where panel (a) puts the estimate on the diagonal, and fails at $r = 6$, where panel (a) puts it at little more than half the truth. The ratio therefore locates the failure inside a single system, using nothing the ground truth supplied.

The Theiler exclusion is the reason the estimator errs most on a transient. With the exclusion removed, it returns at every rank, seed and observer the value maximum likelihood gives on a uniformly sampled curve, while the fast recurrent case is unchanged (Appendix I.2). The value near 29 is a divergent quantity evaluated at the cap on the exclusion; we leave it unclamped (Appendix B.1).

Every system evaluated above is externally driven, whereas training is not, so the evaluation establishes only a conditional result. We therefore ask whether the full-batch runs of section 7 reach the edge of stability [Cohen et al., 2021], where an undriven run might still recur. They do reach it, but oscillation is not recurrence. Some of the runs at the edge never return close to a state already passed through; the diagnostics reject them all alike (Appendix K).

6.2 The delay lag

The delay window spans $(E_{\max} - 1)\tau$ samples. It must cover an appreciable fraction of the oscillation period. Below that the delay coordinates are nearly collinear and the torus never unfolds [Casdagli et al., 1991]. Varying τ alone on a fixed torus moves the estimate over more than an order of magnitude, through the truth only where the window covers a few tenths of a period (figure 5). The lag must therefore be set from the oscillation period. A constructed system supplies that period and a training log does not. Since we cannot recover it from the log, section 7 reads changes and never a level.

6.3 Nuisance factors

At $r = 4$ we vary factors that leave r unchanged over five seeds and twelve observers, measuring each shift against the shift a real change of four components produces (table 9). The shifts worth reporting come from the ramped observer gain and the ramped state amplitude. Smoothness is not among them: shifting the drive band by an octave changes the roughness sharply while moving the estimate by well under a fifth of that reference shift, on every observer. A ramped gain moves the estimate by about a component, so a run whose parameter norm grows steadily will drift by that much at fixed r . This is the confound that matters most in section 7.

Roughness tracks the estimate closely here and again in section 7.3, closely enough to raise a prior question: whether the estimate and the roughness can be separated at all. Appendix M separates them both ways, holding one fixed while the other moves.

6.4 The estimand is a dimension, not an effective rank

Every construction in section 5 drives its r phases at equal amplitude, which makes the active dimension and the effective rank of section 3.1 numerically equal. Where two candidate estimands agree, an experiment cannot say which of them the estimator reports. Driving phase l at amplitude q^l separates them: the path still lies on an r -dimensional torus, while the effective rank falls as q falls. The estimate follows the active dimension (figure 6), staying closer to r than to the effective rank and registering only the phases it can resolve.

The diagnostics. Across this grid the identifiability ratio again separates the accurate cases from the inaccurate ones without overlap, now within a single regime as well as between regimes. The ratio and the trend-crossing count are calibrated on these cases alone, so the band drawn in section 7 records where the accurate cases fell and carries no authority beyond them.

7 Application to grokking

We use two grokking settings. A transformer is trained with weight decay and AdamW on mini-batches, which places it in the stochastic regime of section 3.3; a two-layer perceptron is trained by unregularised full-batch gradient descent (section 2; Appendix J lists every run). In both, the diagnostics reject the level the logs report, so we store the trajectory and measure it directly, then return to the logs to ask whether a change within one run can still be read.

7.1 The estimate may not be read as a dimension in either setting

We run the frozen configuration of section 5.2 on these logs, shortening the window and stride to fit the records and changing nothing else (table 12, figure 7). The transformer runs and the perceptron runs are rejected for different reasons. In neither case is the reason the regime label we would have assigned in advance. The identifiability ratio of the transformer runs lies in a band the deterministic regimes never reach. Their estimates also scatter widely across seeds of one configuration, exactly as figure 1 shows mini-batch noise behaving. The perceptron runs carry the signature of a transient instead: two trend crossings in a whole record, at a lower ratio that is still far outside the band. The value may therefore be read as a dimension in neither setting. A change within one run might still be read, since a fixed configuration adds the same offset to every window of that run, but not at this window, which spans tens of thousands of optimiser steps against a transition a few hundred steps wide. The window is shortened in section 7.3 (Appendix H.3).

Label-matched pairs. The perceptron targets come in pairs differing by a single monomial, both members of which fit the training set completely under identical optimisation. Since the members of a pair remain different polynomials, a separation between them could still be a separation between tasks. The estimate does separate them, but we rest nothing on it. Every one of these runs is transient; the raw loss curves already show the separation (Appendices F, H.1 and H.2).




Figure 2: Simplification of the trajectory at generalisation. Three regularised seeds of modular addition and one S_5 composition, against the same two configurations trained without weight decay, aligned on the generalisation step of the run each matches. Thick line, the mean of the four generalising runs; band, the spread between the two hash families (Appendix G.1). Panels (a) and (b) are the one run whose training curve is stored. Panels (c)–(e) plot PR^{det} and the displacement over windows of 600 steps, 300 for S_5 ; (f) the estimate of section 7.3 on the log of (b).

7.2 The trajectory simplifies at generalisation

The effective rank of section 3.1 is defined whether or not the path returns on itself, so a stored trajectory allows a measurement that a log does not. We store the trajectory of a transformer trained to grok, compressed at every logged step by a random projection of the parameter vector and of the probe logits (Appendix G.1). Near generalisation the effective rank collapses in both spaces, reaching near unity in parameter space. In each space the minimum falls within a few hundred steps of the transition. Both then re-expand (figure 2, table 14). Two competing readings can be set aside. The collapse might follow training progress, but it occurs at generalisation in seeds whose generalisation steps differ several-fold and in a second task. It might record a trajectory that has come to a halt, but the total displacement reaches its floor thousands of steps later and never recovers, whereas the effective rank does recover. One caveat stands. The effective rank measures variation inside the window. Only that variation collapses; d_{act} itself may be unchanged.

The runs without weight decay. The same two configurations trained without weight decay memorise and stay at chance accuracy. Within the transition window they decline much less than the generalising runs, but somewhere in training one of them declines just as far (table 14). Depth therefore does not separate the generalising runs from the memorising ones; the timing of the decline separates them, as does whether it reverses. A low effective rank is reached both by a run that has come to rest and by one about to generalise.

7.3 The log estimate falls where the trajectory does

The test section 7.1 could not perform becomes possible once the window is set in units of the transition, not as a fraction of the record. We re-run the estimator on the same logs, at the window length and midpoints the direct measurement uses, so that both statistics are sampled at the same instants. The frozen configuration does not fit a window that short, while the frozen-configuration requirement forbids choosing a replacement by how well it performs. We therefore report every cell of a grid, taking as headline the cell named by a rule blind to the outcome (Appendix H.3).

On the parameter norm the estimate halves at the generalisation step in all four generalising runs, its minimum falling on the same side of the transition as the minimum of the direct measurement (figure 2f). That drop is among the largest any of the four ever shows, while the other two runs show nothing remarkable there; all but two of the usable grid cells separate the generalising runs from the

other two (table 16). The drop also survives a check on the shape of the log. We compare each run against randomised series built from it that keep its slow shape and its fluctuation spectrum while destroying its fine structure. The fraction of those series falling as steeply as the run itself is smaller for every generalising run than for either of the other two (Appendix H.3).

The same event is thus recorded twice: once by storing the trajectory, once in a log a run writes anyway. The second record is the weaker. Its level cannot be read as a dimension: only the change within a run is used. Its observer is the parameter norm, for which section 6.3 leaves a smoothness explanation open, so the comparison against randomised series carries the claim alone. The roughness of these windows also collapses by an order of magnitude where the estimate merely halves; a simpler statistic registers the change more strongly.

8 Conclusion

In a deterministic recurrent system with a matched delay lag, one scalar log of a suitable observer counts the independent components of the set the parameters occupy, to within about one component and up to about eight. Outside that regime it does not: a transient supplies no returns; mini-batch noise leaves no invariant set to measure. The estimator is usable only alongside the identifiability ratio and the trend-crossing count, which place both grokking settings outside the regime it needs. Our result there comes from the stored trajectory, whose effective rank collapses at generalisation and then re-expands.

Limitations. The measurement needs a recurrent regime and a lag matched to a period no log reveals. Above about eight components it is only ordinal. The collapse comes from four runs against two, distinguished by weight decay, which also drives the parameter norm; it marks the transition only once the transition has happened. It does not reproduce in the full-batch setting. The drop at the matched window is close to the drop a ramped gain alone produces. The edge-of-stability evidence comes from one task (Appendices B.3, H.1, H.3 and K).

References

- Armen Aghajanyan, Sonal Gupta, and Luke Zettlemoyer. Intrinsic dimensionality explains the effectiveness of language model fine-tuning. In *Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 7319–7328, 2021.
- Alessio Ansuini, Alessandro Laio, Jakob H. Macke, and Davide Zoccolan. Intrinsic dimension of data representations in deep neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- Boaz Barak, Benjamin L. Edelman, Surbhi Goel, Sham Kakade, Eran Malach, and Cyril Zhang. Hidden progress in deep learning: SGD learns parities near the computational limit. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- David S Broomhead and Gregory P King. Extracting qualitative dynamics from experimental data. *Physica D: Nonlinear Phenomena*, 20(2-3):217–236, 1986.
- Francesco Camastra and Antonino Staiano. Intrinsic dimension estimation: Advances and open problems. *Information Sciences*, 328:26–41, 2016. doi: 10.1016/j.ins.2015.08.029.
- Liangyue Cao. Practical method for determining the minimum embedding dimension of a scalar time series. *Physica D: Nonlinear Phenomena*, 110(1-2):43–50, 1997. doi: 10.1016/S0167-2789(97)00118-8.
- Martin Casdagli, Stephen Eubank, J. Doyne Farmer, and John Gibson. State space reconstruction in the presence of noise. *Physica D: Nonlinear Phenomena*, 51(1-3):52–98, 1991. doi: 10.1016/0167-2789(91)90222-U.
- Moses Charikar, Kevin Chen, and Martin Farach-Colton. Finding frequent items in data streams. In *International Colloquium on Automata, Languages, and Programming*, pages 693–703, 2002.

380 Jeremy M Cohen, Simran Kaur, Yuanzhi Li, J Zico Kolter, and Ameet Talwalkar. Gradient descent on
381 neural networks typically occurs at the edge of stability. In *International Conference on Learning*
382 *Representations (ICLR)*, 2021. URL <https://openreview.net/forum?id=jh-rTtvkGeM>.

383 Darshil Doshi, Tianyu He, Aritra Das, and Andrey Gromov. Grokking modular polynomials. *arXiv*
384 *preprint arXiv:2406.03495*, 2024.

385 Jean-Pierre Eckmann and David Ruelle. Fundamental limitations for estimating dimensions and
386 lyapunov exponents in dynamical systems. *Physica D: Nonlinear Phenomena*, 56(2-3):185–187,
387 1992. doi: 10.1016/0167-2789(92)90023-G.

388 Elena Facco, Maria d’Errico, Alex Rodriguez, and Alessandro Laio. Estimating the intrinsic di-
389 mension of datasets by a minimal neighborhood information. *Scientific Reports*, 7(1):12140,
390 2017.

391 Andrew M Fraser and Harry L Swinney. Independent coordinates for strange attractors from mutual
392 information. *Physical review A*, 33(2):1134, 1986. doi: 10.1103/PhysRevA.33.1134.

393 Peiran Gao, Eric Trautmann, Byron M. Yu, Gopal Santhanam, Stephen Ryu, Krishna Shenoy, and
394 Surya Ganguli. A theory of multineuronal dimensionality, dynamics and measurement. *bioRxiv*
395 214262, 2017. doi: 10.1101/214262.

396 Peter Grassberger and Itamar Procaccia. Measuring the strangeness of strange attractors. *Physica D:*
397 *Nonlinear Phenomena*, 9(1-2):189–208, 1983.

398 Andrey Gromov. Grokking modular arithmetic. *arXiv preprint arXiv:2301.02679*, 2023.

399 Guy Gur-Ari, Daniel A. Roberts, and Ethan Dyer. Gradient descent happens in a tiny subspace. *arXiv*
400 *preprint arXiv:1812.04754*, 2018.

401 Holger Kantz and Thomas Schreiber. *Nonlinear Time Series Analysis*. Cambridge university press,
402 2004.

403 Matthew B Kennel, Reggie Brown, and Henry DI Abarbanel. Determining embedding dimension for
404 phase-space reconstruction using a geometrical construction. *Physical review A*, 45(6):3403, 1992.
405 doi: 10.1103/PhysRevA.45.3403.

406 Tanishq Kumar, Blake Bordelon, Samuel J. Gershman, and Cengiz Pehlevan. Grokking as the transi-
407 tion from lazy to rich training dynamics. In *International Conference on Learning Representations*
408 *(ICLR)*, 2024. URL <https://openreview.net/forum?id=vt5mnLVIVo>.

409 Elizaveta Levina and Peter J. Bickel. Maximum likelihood estimation of intrinsic di-
410 mension. In *Advances in Neural Information Processing Systems (NeurIPS)*, vol-
411 ume 17, 2004. URL [https://proceedings.neurips.cc/paper/2004/file/](https://proceedings.neurips.cc/paper/2004/file/74934548253bcab8490ebd74afecebb5-Paper.pdf)
412 [74934548253bcab8490ebd74afecebb5-Paper.pdf](https://proceedings.neurips.cc/paper/2004/file/74934548253bcab8490ebd74afecebb5-Paper.pdf).

413 Chunyuan Li, Heerad Farkhoor, Rosanne Liu, and Jason Yosinski. Measuring the intrinsic dimension
414 of objective landscapes. In *International Conference on Learning Representations (ICLR)*, 2018.
415 URL <https://openreview.net/forum?id=ryup8-WCW>.

416 Ziming Liu, Ouail Kitouni, Niklas Nolte, Eric J. Michaud, Max Tegmark, and Mike Williams.
417 Towards understanding grokking: An effective theory of representation learning. In *Advances in*
418 *Neural Information Processing Systems (NeurIPS)*, 2022.

419 Ziming Liu, Eric J. Michaud, and Max Tegmark. Omnigrok: Grokking beyond algorithmic
420 data. In *International Conference on Learning Representations (ICLR)*, 2023. URL <https://openreview.net/forum?id=zDiHoIwa0q1>.

422 David J C MacKay and Zoubin Ghahramani. Comments on ‘maximum likelihood estimation of
423 intrinsic dimension’ by e. levina and p. bickel, 2005. URL [http://www.inference.org.uk/](http://www.inference.org.uk/mackay/dimension/)
424 [mackay/dimension/](http://www.inference.org.uk/mackay/dimension/). Unpublished note.

425 Stephan Mandt, Matthew D Hoffman, and David M Blei. Stochastic gradient descent as approximate
426 Bayesian inference. *Journal of Machine Learning Research*, 18(1):4873–4907, 2017.

427 Sam McCandlish, Jared Kaplan, Dario Amodei, and OpenAI Dota Team. An empirical model of
428 large-batch training. *arXiv preprint arXiv:1812.06162*, 2018.

429 William Merrill, Nikolaos Tsilivis, and Aman Shukla. A tale of two circuits: grokking as competition
430 of sparse and dense subnetworks. *arXiv preprint arXiv:2303.11873*, 2023.

431 Neel Nanda, Lawrence Chan, Tom Lieberum, Jess Smith, and Jacob Steinhardt. Progress measures for
432 grokking via mechanistic interpretability. In *International Conference on Learning Representations*
433 (ICLR), 2023.

434 Pascal Tikeng Notsawo Jr., Hattie Zhou, Mohammad Pezeshki, Irina Rish, and Guillaume Dumas.
435 Predicting grokking long before it happens: A look into the loss landscape of models which grok.
436 *arXiv preprint arXiv:2306.13253*, 2023.

437 A. R. Osborne and A. Provenzale. Finite correlation dimension for stochastic systems with power-law
438 spectra. *Physica D: Nonlinear Phenomena*, 35(3):357–381, 1989.

439 Phil Pope, Chen Zhu, Ahmed Abdelkader, Micah Goldblum, and Tom Goldstein. The intrinsic dimen-
440 sion of images and its impact on learning. In *International Conference on Learning Representations*
441 (ICLR), 2021. URL <https://openreview.net/forum?id=XJk19XzGq2J>.

442 Alethea Power, Yuri Burda, Harri Edwards, Igor Babuschkin, and Vedant Misra. Grokking: Gen-
443 eralization beyond overfitting on small algorithmic datasets. *arXiv preprint arXiv:2201.02177*,
444 2022.

445 A. Provenzale, L. A. Smith, R. Vio, and G. Murante. Distinguishing between low-dimensional
446 dynamics and randomness in measured time series. *Physica D: Nonlinear Phenomena*, 58(1-4):
447 31–49, 1992.

448 Olivier Roy and Martin Vetterli. The effective rank: a measure of effective dimensionality. In
449 *European Signal Processing Conference (EUSIPCO)*, pages 606–610, 2007.

450 Levent Sagun, Utku Evci, V. Uğur Güney, Yann Dauphin, and Léon Bottou. Empirical analysis of the
451 Hessian of over-parametrized neural networks. *arXiv preprint arXiv:1706.04454*, 2017.

452 Tim Sauer, James A. Yorke, and Martin Casdagli. Embedology. *Journal of Statistical Physics*, 65
453 (3-4):579–616, 1991.

454 Thomas Schreiber and Andreas Schmitz. Improved surrogate data for nonlinearity tests. *Physical*
455 *Review Letters*, 77(4):635–638, 1996.

456 Jaroslav Stark. Delay embeddings for forced systems. I. Deterministic forcing. *Journal of Nonlinear*
457 *Science*, 9(3):255–332, 1999. doi: 10.1007/s003329900072.

458 Jaroslav Stark, David S Broomhead, ME Davies, and J Huke. Delay embeddings for forced
459 systems. II. Stochastic forcing. *Journal of Nonlinear Science*, 13(6):519–577, 2003. doi:
460 10.1007/s00332-003-0534-4.

461 George Sugihara and Robert M May. Nonlinear forecasting as a way of distinguishing chaos from
462 measurement error in time series. *Nature*, 344(6268):734–741, 1990. doi: 10.1038/344734a0.

463 George Sugihara, Robert May, Hao Ye, Chih-hao Hsieh, Ethan Deyle, Michael Fogarty, and Stephan
464 Munch. Detecting causality in complex ecosystems. *Science*, 338(6106):496–500, 2012. doi:
465 10.1126/science.1227079.

466 Floris Takens. Detecting strange attractors in turbulence. In *Dynamical systems and turbulence*,
467 *Warwick 1980*, pages 366–381. Springer, 1981. doi: 10.1007/BFb0091924.

468 James Theiler. Spurious dimension from correlation algorithms applied to limited time-series data.
469 *Physical Review A*, 34(3):2427–2432, 1986.

470 James Theiler. Some comments on the correlation dimension of $1/f^\alpha$ noise. *Physics Letters A*, 155
471 (8-9):480–493, 1991.

- James Theiler, Stephen Eubank, André Longtin, Bryan Galdrikian, and J. Doyne Farmer. Testing for nonlinearity in time series: the method of surrogate data. *Physica D: Nonlinear Phenomena*, 58 (1-4):77–94, 1992.
- Vimal Thilak, Etai Littwin, Shuangfei Zhai, Omid Saremi, Roni Paiss, and Joshua M Susskind. The slingshot mechanism: An empirical study of adaptive optimizers and the loss landscape. *arXiv preprint arXiv:2206.04817*, 2022. URL <https://arxiv.org/abs/2206.04817>.
- Vikrant Varma, Rohin Shah, Zachary Kenton, János Kramár, and Ramana Kumar. Explaining grokking through circuit efficiency. *arXiv preprint arXiv:2309.02390*, 2023.
- Kilian Weinberger, Anirban Dasgupta, John Langford, Alex Smola, and Josh Attenberg. Feature hashing for large scale multitask learning. In *International Conference on Machine Learning (ICML)*, pages 1113–1120, 2009.

Appendix outline

These appendices are supplementary. The body stands without them; each holds the data, the construction or the audit behind a claim made there. The paragraphs below say which.

Appendix A states, without proofs, the four published results the estimator is assembled from: the delay embedding theorem of Takens [1981], its extension to fractal sets, the version for systems driven by an external input on which section 3.3 rests its treatment of mini-batch noise, and the maximum likelihood estimator of an intrinsic dimension. It closes with the limits a finite record places on any such statistic.

Appendix B records the procedure as run: the estimator as an algorithm, with its constants and the points at which the implementation departs from equation (4); the scalar observers of section 5.2; and the configurations at which the estimator was frozen, with the grids they were selected on and the qualifications those choices carry.

Appendix C scores section 5 in full: the summary table the body cites, the result for each observer separately, and every simpler statistic we could compute on the same data, set beside the estimate. Three of those comparisons go against the estimator and are reported as such.

Appendix D gives the three experiments section 6 summarises, on the delay lag, on the nuisance factors and on unequal drive amplitudes. It adds a fourth that the body does not report: a test that a real change in the active dimension is detected when one occurs.

Appendix E describes how the image-data system of section 5.2 is built in the three variants the paper scores and shows by measurement that each excites as many independent directions as intended. Estimates on that system are scored against the measured values, never against the nominal ranks.

Appendix F gives the diagnostic values behind section 7.1 for every training log and places each log in the plane of the identifiability ratio and the trend-crossing count. It also shows the label-matched pairs of section 7.1 beside their raw loss curves, which limit the reading that can be given to the separation between the members of a pair.

Appendix G covers the measurement of section 7.2: the random projection that makes storing a whole trajectory affordable, the checks that validate it, the spectral guarantee those checks do not supply, and the collapse reported run by run.

Appendix H makes one argument three times. A statistic computed over a window of a given length cannot localise anything shorter than that window, so a feature that fails to appear may be beyond the resolution of the measurement rather than absent from the run. A decline that does appear may still not be specific to generalisation. The three parts make that argument for the effective rank of the trajectory in the full-batch setting, for a run whose estimate declines although it never generalises, and for the shortened window of section 7.3, where the qualifications on that result are also stated.

Appendix I audits the Theiler exclusion, the one estimator setting that determines the value returned on a transient and also the setting at which this paper departs from its own protocol. It measures the cost of that departure and varies the exclusion from zero upwards to show that the dimension of a transient is one and is recovered only when the exclusion is absent.

521 **Appendix J** lists every training run the paper uses, with its task, architecture, length and outcome. It
 522 also states where our two grokking settings differ from the published configurations they are taken
 523 from: mini-batch sampling in the transformer setting and plain gradient descent in the perceptron
 524 setting place the runs in the regimes section 7.1 assigns them.

525 **Appendix K** asks whether an undriven system can be deterministic and recurrent at once, as the
 526 estimate requires. It tests gradient descent at the edge of stability and shows besides that writing the
 527 log less often can remove the oscillation from the record altogether.

528 **Appendix L** varies the embedding dimension and the record length independently. It separates the
 529 level the estimator returns at high rank from the highest rank at which it still tracks the truth, then
 530 reports that neither candidate explanation of the ceiling survives.

531 **Appendix M** establishes that the estimate and the roughness of a log can be separated at all, by
 532 holding each fixed in turn while the other moves. It also reports two findings that count against the
 533 estimator.

534 **Appendix N** states the compute used, the storage and time a compressed trajectory costs, and the
 535 material to be released.

536 A The results behind the estimator

537 The procedure of section 3.2 is assembled from four published results, stated here without proof and
 538 in the form the paper uses. Every one of the conditions they impose is an assumption about the run
 539 rather than a property we can verify from the log.

540 Throughout, f is the map advancing the system by one step, ϕ the scalar observation, and

$$\Phi_{f,\phi}(s) = (\phi(s), \phi(f(s)), \dots, \phi(f^{E-1}(s))) \in \mathbb{R}^E \quad (5)$$

541 the delay map.

542 **Theorem 1** (Delay embedding; [Takens, 1981](#)). *Let M be a compact manifold of dimension d and let*
 543 *$E > 2d$. For a generic pair (f, ϕ) with $f \in C^2(M, M)$ a diffeomorphism and $\phi \in C^2(M, \mathbb{R})$, the*
 544 *delay map equation (5) is an embedding of M into $\mathbb{R}^E$.*

545 An embedding is one-to-one and an immersion, so distances in $\mathbb{R}^E$ between reconstructed points
 546 are a faithful, though distorted, image of distances between states. The set they fill is the object the
 547 neighbour statistic below measures (figure 3). Theorem 1 requires a manifold.

548 The sets a training run occupies need not be manifolds; the extension below is the version the paper
 549 relies on.




Figure 3: The set the neighbour search measures, in the plane of the first two delay coordinates. One window of the image-data system of section 5.2 under each excitation, reconstructed at the frozen configuration. A single phase closes an orbit, two fill a surface, a converging run never returns, and mini-batch noise fills the space it is given.

550 **Theorem 2** (Fractal delay embedding prevalence; [Sauer et al., 1991](#)). *Let $A \subset \mathbb{R}^n$ be compact*
 551 *with box-counting dimension d , invariant under a smooth diffeomorphism f defined on an open*
 552 *neighbourhood of A , and let $E > 2d$ be an integer. Assume that for each $p \leq E$ the set of periodic*
 553 *points of f of period p has box-counting dimension less than $p/2$, and that the linearisation of f at*

each such orbit has distinct eigenvalues. Then for almost every smooth ϕ , in the sense of prevalence, the delay map equation (5) is one-to-one on A and an immersion on every compact subset of a smooth manifold contained in A .

The dimension in the condition $E > 2d$ is the box-counting dimension of equation (1), which need not be an integer, so the criterion applies to the sets a training run occupies and not only to tori. Prevalence, in turn, is a measure-theoretic notion of “almost every” on an infinite-dimensional space. Mini-batch sampling makes the update depend on a draw from the data as well as on the state, so neither theorem above applies as stated.

Theorem 3 (Delay embedding for forced systems; Stark, 1999, Stark et al., 2003). *Let $s_{t+1} = f(s_t, w_t)$ with the drive w_t generated by a measure-preserving system on a space Ω , either deterministically or as an independent draw selecting among finitely many maps. Under genericity conditions on f and ϕ analogous to those of theorem 1, and for E exceeding twice the dimension of a fibre, the map sending a state to its delay vector equation (5) is an embedding of each fibre M_ω , that is, of the set of states compatible with one realisation ω of the drive.*

The guarantee holds fibre by fibre, whereas the sample available in practice is one trajectory, which lies in a different fibre at every step. No invariant measure on a fibre is being sampled, so the neighbour statistic below estimates nothing.

Theorem 4 (Maximum likelihood intrinsic dimension; Levina and Bickel, 2004, MacKay and Ghahramani, 2005). *Let points be drawn from a smooth density supported on a d -dimensional manifold embedded in $\mathbb{R}^E$. Model the points falling within a small ball about $\mathbf{y}_i$ as a homogeneous Poisson process in the d -dimensional radial variable, with the density constant over the ball. Writing $r_j^{(i)}$ for the distance from $\mathbf{y}_i$ to its j -th nearest neighbour, the maximum likelihood estimate of d from the first m neighbours of that point is*

$$\hat{d}_m(\mathbf{y}_i)^{-1} = \frac{1}{m-1} \sum_{j=1}^{m-1} \log \frac{r_m^{(i)}}{r_j^{(i)}}. \quad (6)$$

Under this model the reciprocal is unbiased and the estimate itself is not.

The estimator is therefore a local statistic: it reports the dimension at the scale set by the m nearest neighbours, the finite ε of equation (1) rather than a limit as $\varepsilon \rightarrow 0$. A record of N points resolves a dimension of at most about $2 \log_{10} N$ under the argument of Eckmann and Ruelle [1992]. A series with a power-law spectrum returns a finite, reproducible and entirely spurious value fixed by its spectral exponent [Osborne and Provenzale, 1989, Theiler, 1991]. Neither bound explains the ceiling we measure, which appendix L takes up.

B The estimator, its configuration and the observers

B.1 The estimator

Algorithm 1 is the estimator as it was run. It computes equation (4) with three departures that affect the values reported.

Algorithm 1 Pooled Levina–Bickel intrinsic dimension of a delay reconstruction

Require: series $x_{1:n}$; lag τ ; embedding dimension E ; neighbours m ; Theiler exclusion W_T ; dither

σ
1: $z_t \leftarrow (x_t - \text{mean } x) / \text{std } x + \sigma \epsilon_t, \quad \epsilon_t \sim \mathcal{N}(0, 1)$ ▷ breaks exact ties
2: $\mathbf{y}_t \leftarrow (z_t, z_{t-\tau}, \dots, z_{t-(E-1)\tau})$ for $t = (E-1)\tau + 1, \dots, n$
3: $S \leftarrow 0; N \leftarrow 0; n_{\text{dist}} \leftarrow 0; n_{\text{ratio}} \leftarrow 0$
4: **for each** $\mathbf{y}_i$ **do**
5: $r_1^{(i)} \leq \dots \leq r_m^{(i)} \leftarrow$ distances to the m nearest $\mathbf{y}_j$ with $|i - j| > W_T$
6: $r_j^{(i)} \leftarrow \max(r_j^{(i)}, \varepsilon_{\text{dist}})$ for each j ; count those clipped into n_{dist}
7: $S_i \leftarrow \max(\sum_{j=1}^{m-1} \log(r_m^{(i)} / r_j^{(i)}), \varepsilon_{\text{ratio}})$; count clips into n_{ratio}
8: $S \leftarrow S + S_i; \quad N \leftarrow N + 1$
9: **end for**
10: mark the window degenerate if n_{dist} or n_{ratio} exceeds 1 % of its denominator
11: **return** $(N(m-1) - 1) / S$

588 *The returned value is $(N(m-1) - 1) / S$, the exactly unbiased pooled form, which differs from*
589 *equation (4) by less than 10^{-5} at the sizes used here. A window marked degenerate is returned rather*
590 *than discarded. No published cell of table 12 contains a marked window. The returned value is never*
591 *clamped to E , so a divergent estimate is reported as divergent.*

592 The Theiler exclusion is $W_T = \max((E-1)\tau, t_{\text{acf}})$, where t_{acf} is the first lag at which the
593 autocorrelation falls below $1/e$, capped at 150 samples because the neighbour query cost grows
594 linearly with it; appendix I.1 measures the effect of that cap. The floors $\varepsilon_{\text{dist}}, \varepsilon_{\text{ratio}}$ and the dither σ
595 are $10^{-8}, 10^{-5}$ and 10^{-9} .

596 B.2 The observers

597 Twelve scalar observers are recorded on the system of section 5.2, in five families; table 1 defines
598 each.

Table 1: The scalar observers. $\mathbf{c}$ is the coordinate in the k -dimensional subspace and $\boldsymbol{\theta}_0$ the frozen head at initialisation; $\mathbf{L}$ is the matrix of probe logits, where the probe is a fixed held-out set drawn once before training; $\mathbf{g} = \partial \ell / \partial \mathbf{c}$. The projection directions $\mathbf{u}, \mathbf{v}$ and $\mathbf{a}$, and the rotation $\mathbf{R}$, are drawn once from a fixed seed and never redrawn.

name	family	definition
instantaneous loss	loss	the loss on the current minibatch
full-batch loss	loss	the loss on the whole training set
probe loss	loss	cross-entropy on the probe set
parameter norm	norm	$(\ \boldsymbol{\theta}_0\ ^2 + \ \mathbf{c}\ ^2)^{1/2}$
subspace norm	norm	$\ \mathbf{c}\ _2$
function-space norm	norm	$\ \mathbf{L}\ _F$
gradient norm	gradient	$\ \mathbf{g}\ _2$
gradient projection	gradient	$\langle \mathbf{u}, \mathbf{g} \rangle$
fixed parameter projection	projection	$\langle \mathbf{v}, \mathbf{R}\mathbf{c} \rangle$
function-space projection	function	$\langle \mathbf{a}, \text{vec } \mathbf{L} \rangle$
margin	function	mean probe logit gap, true class less best other
probe accuracy	function	accuracy on the probe set

599 B.3 Frozen configurations

600 The estimator has one frozen configuration for each range of the active dimension. Each was selected
601 once, on data withheld from every later experiment (table 2).

Table 2: The frozen configurations and their selection. Only the swept parameters were varied; the remainder were fixed before any estimate was computed. Settings common to both configurations are not repeated: $m = 20$ neighbours, a Theiler exclusion of at least the embedding span (appendix B.1). The neighbour count and the exclusion were swept in the eight-component grid, in which the two Theiler rules return bit-identical values, so the selection between them is arbitrary.

	eight-component	twenty-component
$E_{\max}$	20	40
τ	4	16
window / stride	8 000 / 2 000	8 000 / 4 000
swept	$E_{\max}, \tau, m, \text{Theiler, window}$	$E_{\max}, \tau$
grid size	48	4
selection seeds	90, 91, 92	0, 1, 2
selection ranks	2, 4, 6	2, 6, 10, 14, 18
withheld seeds	0–3	none
withheld ranks	1, 3, 5, 8	the other fifteen
selection error, best	0.17	2.47
selection error, worst	2.03	3.35

602 The selection objective in each case is the error of the median estimate against the median measured
603 effective rank, over the four observers reserved for selection. The eight-component objective adds a
604 penalty on the rank correlation and on the spread across seeds.

605 *Each configuration lies at a grid boundary.* The eight-component configuration is at the top of all five
606 swept axes, the twenty-component one at the top of both its axes. A wider grid could move the level.

607 *The selection error spans a factor of twelve on identical data* (table 2), so the level is as much a
608 choice as a measurement. Only the withheld numbers of section 5.2 may be read as recovery.

609 *The stride varies between experiments.* Table 3 states the window geometry each experiment uses;
610 none of them overrides a setting of the estimator. The constructed systems produced the 0.90 of
611 section 5.2, so “at the frozen configuration” is true there of the estimator and not of the stride.

Table 3: The window and stride of each experiment, with their departures from the frozen configuration of table 2. Lengths are in samples of the log; the second stride applies to the longer records.

experiment	window	stride
frozen configuration	8 000	2 000
constructed systems (section 5)	8 000	3 000, or 3 666
training logs (section 7)	a third of the record	1 000

612 C Per-observer results and alternative statistics

613 Table 4 gives the evaluation of section 5 in full.

Table 4: How accurately the estimator recovers the active dimension on each system, with the requirement of section 4 that each fails. Systems are ordered by increasing realism. The unit of analysis is one (rank, seed, observer) run, each side of the error being the median over that run’s sliding windows; ρ is over ranks after a median over seeds. The oscillating matrix has no optimiser, so the zero-learning-rate check does not apply to it. The function-subspace row fails constructed truth at its three highest ranks, which reach an effective rank of about $0.86r$ against the $0.9r$ the check asks.

system	truth	range	MAE	ρ	requirement failed
oscillating matrix	constructed	1–8	0.27	1.00	<i>not applicable</i>
online linear regression	constructed	1–20	0.79	0.99	frozen configuration
logistic regression	constructed	1–20	0.81	1.00	frozen configuration
frozen nonlinear decoder	constructed	1–20	1.58	0.94	zero learning rate
perceptron in a k -subspace	constructed	1–20	0.92	0.99	zero learning rate
image data, parameter subspace	measured	1–8	0.90	1.00	none
image data, parameter subspace	measured	1–20	1.75	1.00	withheld ranks
image data, function subspace	measured	1–8	0.28	1.00	all of them

614 The function-subspace row is marked as failing every requirement, but the zero-learning-rate check
615 was never re-run on it.

616 All numbers in this appendix are computed on the withheld ranks $r \in \{1, 3, 5, 8\}$ and the withheld
617 seeds of section 5.2, in the deterministic recurrent regime, scored against the measured effective rank
618 rather than against r . Where an aggregation over seeds is needed, the median over the four seeds at
619 each rank is taken before scoring. Table 5 reconciles the three errors this paper quotes for that one
620 measurement.

Table 5: Why three different errors are quoted for one measurement. Each row is the same eight-component recovery of section 5.2, on withheld ranks and withheld seeds in the deterministic recurrent regime, summarised in a different order. The order is stated wherever one of these figures appears.

aggregation	reported in	error
mean over (rank, seed, observer) runs	section 5.2, table 4	0.90
median over observers of each observer’s error	table 6	0.48
median over seeds and observers at each rank, then scored	table 8	0.382

Table 6: The accuracy of each scalar observer, against the smoothness that might explain it. One row per observer, except that the parameter norm and the subspace norm share one, their errors agreeing to two decimals. The rank correlation is 1.00 for every observer, so it has no column. The last column records whether the roughness orders the withheld ranks perfectly.

observer	family	MAE	slope	roughness orders rank
probe loss	loss	0.60	0.82	no
function-space norm	norm	0.39	0.96	no
margin	function	0.46	0.83	no
parameter, subspace norm	norm	0.41	1.04	no
fixed parameter projection	projection	0.34	0.83	no
gradient projection	gradient	0.70	0.72	no
function-space projection	function	0.51	0.78	no
gradient norm	gradient	1.81	1.58	no
full-batch loss	loss	1.87	1.51	no
instantaneous loss	loss	0.61	0.73	no
probe accuracy	function	degenerate in every window		

621 Two observers are excluded from every aggregate. Probe accuracy is quantised, which collapses the
622 delay vectors onto one another and marks every one of its windows degenerate. The instantaneous
623 mini-batch loss fails the zero-learning-rate check: its estimate still climbs by about six components

there, produced entirely by the drive. The other eleven observers are exactly constant at zero learning rate, so no estimate is defined for them.

An aggregate over observers answers a different question from the best row of table 6. Nothing in a training log says which observer is the accurate one, since that judgement needs the ground truth a log does not carry, so a reader can act only on the typical observer, not the best.

For no scored observer does the roughness order the four withheld ranks; the best reaches a rank correlation of 0.8. Figure 4 shows the estimate and the roughness along one record.

Smoothness alone therefore accounts for no observer’s agreement with the truth.




Figure 4: Whether four logs that look alike separate when the estimator reads them. The image-data system of section 5.2 at four ranks and a seed the configuration was not selected on, through the parameter norm. Above, four hundred samples of each log. (a) The estimate over sliding windows, the measured effective rank dotted. (b) The roughness of the same windows.

Table 7: How far recovery extends at $k = 20$. Image data, median over three seeds and over the eight observers recorded there, which reports the typical observer rather than the best one; the truth is the measured effective rank. The last row is the spread of MG across seeds at fixed rank.

true r	1	2	4	6	8	10	12	14	17	20
MG	1.10	2.05	4.51	6.51	7.45	9.92	10.23	11.40	13.90	15.44
PR _{delay}	2.00	2.26	3.79	5.42	7.90	9.45	9.43	11.25	11.43	13.42
spectral PR	1.09	1.27	2.13	2.89	4.10	4.60	5.51	6.22	6.81	7.48
spread across seeds	0.87	0.51	1.35	1.81	2.64	3.59	2.81	2.07	3.41	4.24

Table 8: Whether any simpler statistic recovers the active dimension as well as the estimator does. All are computed on the image-data system of section 5.2. The two ranges are not aggregated alike: at $r \leq 8$ each statistic is scored on the withheld ranks alone, at $r \leq 20$ on all twenty, the selection ranks included, so only the comparison down a column is out of sample. Statistics marked $\dagger$ need a neighbour search; the rest are one-line computations on the series or its spectrum.

statistic	$r \leq 8$		$r \leq 20$	
	MAE	ρ	MAE	ρ
MG $\dagger$	0.382	1.00	1.75	0.997
LB $\dagger$	0.40	1.00	1.52	0.997
TwoNN $\dagger$	2.13	0.80	2.79	0.880
PR _{delay}	0.91	1.00	2.48	0.985
spectral PR, 256 bands	1.25	1.00	5.83	0.998
spectral PR, 1024 bands	1.08	1.00	4.78	0.979
spectral PR, native resolution	0.69	1.00	2.95	0.976
roughness	3.65	0.20	9.89	-0.379

Three qualifications attach to table 8. The uncorrected per-point average LB is more accurate than MG at $r \leq 20$: the correction of section 3.2 removes a positive bias, which on a saturating estimate was partly cancelling the saturation. The spectral participation ratio at 256 bands orders the ranks at $r \leq 20$ more faithfully than MG, at three times the error. The linear delay participation ratio needs no neighbour search at all, yet MG beats it by only a factor of about two at $r \leq 8$.

D Conditions of validity in detail

The delay lag. Figure 5 varies the lag τ on a fixed torus, leaving the system itself unchanged. Below a tenth of a period every rank returns the same value near 2.5 and the ordering is lost; between three tenths and a half of a period the estimate follows the rank; above one period it diverges.




Figure 5: The delay lag, varied on one system: a fixed torus, two seeds, six lags spanning a twentieth of a period to one and a half periods. Marker shape is the true rank, the marker the median of the seeds and the bar their range. (a) The estimate. (b) The same divided by the true rank, so one reference line replaces six; the shaded column is three tenths to a half of a period.

Nuisance factors. Table 9 varies factors that leave the active dimension at four. The row for rotated coordinates is scored on the fixed parameter projection alone, since the other eleven observers are invariant to a rotation by construction; that row measures a property of the observer rather than of the estimator.

Table 9: Shift in the estimate produced by factors that leave the active dimension unchanged, median over five seeds and twelve observers. The first row is the seed-to-seed variability of an unchanged configuration and sets the scale for the others. The last column is a percentage of a real four-component change, which the change-detection experiment below observes as a fall of 3.96. A constant rescaling of the observer moves the estimate by 0.000 and is omitted.

factor	$ \Delta \text{MG} $	as % of a real four-component change
none (baseline)	0.39	10
learning rate halved	0.37	9
drive band up one octave	0.50	13
drive band down one octave	0.78	20
coordinates rotated	0.28	7
observer gain ramped $10\times$	1.53	38
state amplitude ramped $4\times$	1.57	39

645 **Detection of a change.** Recovery of a static level does not establish that a change is detected. We
646 run three consecutive segments $r_{\text{high}} \rightarrow r_{\text{low}} \rightarrow r_{\text{high}}$ with nothing else altered, re-measuring the
647 ground truth on each. In the deterministic recurrent regime a true fall of 4.00 components is observed
648 as 3.96, is detected both downward and upward on over nine tenths of the runs, and recovers to within
649 0.05 of the level before it. Under stochastic excitation the same true change is observed as 0.09.

650 The rule declares a change when the estimate falls below the median of the pre-switch segment
651 by three times the window-to-window scatter of that segment, for two consecutive windows. The
652 threshold uses the pre-switch segment alone.

653 Table 10 gives the three schedules on the parameter norm. The level error is 0.45 at the high level
654 and 0.22 at the low one, so the estimate is more accurate at the lower level.

Table 10: Whether a change of the active dimension is recovered at both levels and after it reverses. Three schedules $r_{\text{high}} \rightarrow r_{\text{low}} \rightarrow r_{\text{high}}$, on the parameter norm. At $6 \rightarrow 2 \rightarrow 6$ the high level is over-reported by a component.

schedule	high	low	high again
$4 \rightarrow 1 \rightarrow 4$	4.06	0.84	4.06
$6 \rightarrow 2 \rightarrow 6$	7.00	2.17	6.95
$8 \rightarrow 3 \rightarrow 8$	8.32	3.34	8.28

655 **Unequal drive amplitudes.** Figure 6 gives the measurement of section 6.4 in full: the system of
656 section 5.1 with phase l driven at amplitude q^l , three seeds, observed through the squared Frobenius
657 norm and estimated with the frozen eight-component configuration. The manifold dimension is r
658 throughout.




Figure 6: Whether the estimate follows the active dimension or the effective rank. Driving phase l at amplitude q^l holds the active dimension at r while the effective rank falls with q . The system of section 5.1, three seeds, one panel per r ; the dotted line is r . Filled circles, the estimate, median over seeds, with standard deviation as bars; open squares, the measured effective rank, which matches its closed form to two decimals and so has no spread.

At $r = 2$ and $r = 4$ the estimate is flat in q while the effective rank declines, as intended. At $r = 6$ and $r = 8$ the estimate declines as well. The amplitudes show why: at $q = 0.5$ and $r = 8$ the weakest phase is driven at 0.008 of the strongest, far below the neighbourhood scale of the estimator, so the torus it resolves really does have fewer dimensions than the one that was built. The estimator therefore reports the number of phases it can resolve. That is neither r nor the effective rank; the three coincide only when the amplitudes are equal.

E System constructions

The image-data system of section 5.2. The data is partitioned into twelve fixed groups, the same partition at every rank so that the partition is not itself a function of r . The loss weights of r of those groups are modulated by r rationally independent sinusoids filling one octave. Modulating group j tilts the gradient along a direction ϕ_j . These directions are neither orthogonal nor of equal effect, since random data groups have correlated gradients, so a mixing matrix $\text{pinv}(\Phi)Q$ with Q an orthonormal basis of $\text{range}(\Phi)$ orthogonalises them. The per-phase scalar response gain of the linearised update is then divided out. Without this correction the trajectory participation ratio falls well below r .

Table 11 verifies that each construction excites as many independent directions as it is meant to. Two of its rows report an outcome rather than a setting. A decaying transient does not recur whatever r is, since a converging trajectory is a curve. Plain mini-batch descent holds the same rank regardless of r , because the rank of its noise is fixed by the per-example gradient covariance of the data; projecting that noise onto r directions recovers only part of the intended rank, so the estimates in that regime are scored against the measured value.

Table 11: Whether each construction excites as many independent directions as intended. Measured trajectory participation ratio by excitation regime and nominal rank, median over the four seeds, for the ten available directions of section 5.2. Every estimate on this system is scored against these values, never against r .

regime	$r = 1$	2	4	6	8	excitation
qp	1.00	2.00	4.00	6.00	8.00	r sinusoids, fast drive
qp_slow	1.00	2.00	4.00	5.99	7.98	the same at a training-log timescale
qp_nopre	1.00	2.00	3.73	5.98	7.77	qp without the preconditioner
mixed	1.00	2.00	4.00	6.00	7.99	qp plus noise in the same directions
noise	1.00	2.00	4.00	5.99	7.99	injected rank- r gradient noise
noise_nopre	1.09	1.98	3.71	5.24	6.75	the same without the preconditioner
batch_proj	1.00	1.74	3.08	4.19	5.38	real mini-batch noise projected to rank r
batch	6.58	6.58	6.56	6.60	6.56	plain mini-batch descent
gd	1.04	1.05	1.04	1.04	1.06	a decaying transient
qp, $\eta = 0$	1.00	1.00	1.00	1.00	1.00	the zero-learning-rate check

The functional rank of the backbone is 10 at every rank, seed and regime, a property of the frozen backbone rather than of the excitation. The drive-direction matrix is perfectly conditioned at $r = 1$ by construction and has a median condition number of 16.2 at $r = 8$, the anisotropy the mixing matrix exists to remove.

The function-space construction. A local adapter is built around a trained network of widths 64, 16 and 10. Its function-space Jacobian is QR-whitened, so that the first r columns have rank r at equal local scale. A quasiperiodic teacher excites every adapter coordinate and the student tracks it by exact backpropagation through both layers. The functional, trajectory-covariance and update-covariance ranks were all verified to equal r .

The system at twenty available directions. The same construction at $k = 20$ reproduces the nominal rank over $r = 1, \dots, 20$ in the recurrent regime. The other regimes behave as they do at $k = 10$: injected noise tracks r , projected mini-batch noise falls short of it, and the decaying transient stays near unity throughout.

F Diagnostics and figures for the grokking application

Table 12 gives the diagnostics behind section 7.1.

Table 12: Whether the estimate may be read as a dimension in either grokking setting. The estimator and its diagnostics on training logs at $E_{\max} = 20$: the regularised stochastic setting through the parameter norm, the unregularised full-batch setting through the training loss. Every cell is a median over the run’s sliding windows. Seven of the seventeen runs of table 19 are shown; figure 7 plots all seventeen. Runs that have not generalised within the allotted steps are censored, not negative.

run	generalises	MG	$\rho_{\text{id}}^{\text{ident}}$	roughness	PR_{delay}	crossings
transformer, full data	yes	13.32	1.53	0.097	1.03	204
transformer, reduced data, s0	no	4.70	1.59	0.067	1.50	161
transformer, reduced data, s2	no	4.42	1.55	0.139	4.04	153
transformer, no weight decay	no	15.60	1.76	< 0.001	1.00	2
perceptron, $n + m$	yes	15.08	1.42	0.001	1.00	2
perceptron, $n \cdot m$	yes	14.97	1.39	0.001	1.00	2
perceptron, $n^3 + nm^2 + m$	no	15.41	1.51	< 0.001	1.00	2

Figure 7 places every training log of section 7 in the plane of the identifiability ratio and the trend-crossing count, against the region in which every accurate case of section 6 fell. No run lies in that region and none is close: every one is unstable in the embedding dimension, at ratios well above

the 1.10 the region allows. On the other axis the transformer and perceptron runs fail in opposite directions, since the perceptron series are monotone.




Figure 7: Whether the estimate may be read as a dimension on any grokking log. Every log of section 7 in the plane of the identifiability ratio and the trend-crossing count, the transformer runs observed through the parameter norm and the perceptron runs through the training loss. Colour is the regime of section 3.3, marker shape the setting. The shaded rectangle is where every accurate case of section 6 fell. It is a record of those cases, not a test.

Figure 8 shows the separation of section 7.1 and the mechanism behind it.




Figure 8: Whether the estimate tells apart two targets that differ by one monomial, only one of them representable by the architecture. (a) The pairs of section 7.1, on the training loss; both members of a pair reach full training accuracy. Marker, median over the sliding windows; thick bar, interquartile range; thin line, full range. All runs are transient, hence one colour; fill and dash mark generalisation. The medians separate every pair, but the ranges overlap, because the estimate drifts upward over training. (b) One pair in full: the representable target drives the loss orders of magnitude lower, so the two runs differ in the shape of the loss curve rather than in any dimension.

G The direct trajectory measurement

G.1 The trajectory sketch

The direct measurement of section 7.2 needs the trajectory itself: the whole parameter vector at each of two to three thousand logged steps. Every logged step is therefore compressed to $\mathbb{R}^{1024}$ by a CountSketch [Charikar et al., 2002, Weinberger et al., 2009], a sparse random projection, drawn twice from two independent hash families so that the two copies can be compared. Two spaces are recorded: the parameters, and the probe logits after centring each example over classes and normalising to unit length. The normalisation is necessary because the raw logit scale is mechanically coupled to weight

709 decay; without it, a signal taken from the logits could not be distinguished from a measurement of
710 weight decay.

711 We do not verify a spectral guarantee for the compression. The two hash families supply a direct
712 estimate of the error instead: every reported value is the mean of the two projections, with the
713 disagreement between them recorded beside it. That disagreement is a per cent of the value at the
714 median and under nine per cent at its worst, smaller than every effect section 7.2 reports.

715 *The sketch is accurate on synthetic trajectories of known rank.* Table 13 passes such trajectories
716 through the sketch and compares the result against the same statistic computed on the uncompressed
717 trajectory. The sketch is accurate to 0.11 components throughout. Nothing above rank ten is tested,
718 however, whereas the plateaus in figure 2 are higher.

719 *It agrees with exact arithmetic on a real run.* On the perceptron architecture, where the parameter
720 count is small enough to allow it, the participation ratio computed on all 145 500 raw parameters
721 agrees with the sketched value to three decimals.

722 *Recording the trajectory does not perturb training.* The log is required to be bit-identical with and
723 without the recording, over four hundred steps at reduced probe size. That check matters because a
724 single global random stream supplies the train-validation split, the initialisation and the mini-batch
725 order in turn, so one stray draw changes the initial weights and can destroy grokking.

Table 13: The cost of the compression. The sketch against the same statistic computed without it, on synthetic trajectories of 60 samples and known rank in 145 500 dimensions. The last column is the cost of the compression. The difference between the first two columns is the cost of a sixty-sample participation ratio and is not attributable to the sketch.

true rank	PR, uncompressed	PR, sketched	sketch – uncompressed
1	1.00	1.00	0.00
2	1.99	1.99	0.00
5	4.56	4.54	−0.02
10	8.44	8.33	−0.11

726 Windows are sixty logged rows: 600 optimiser steps for the modular runs, 300 for S_5 . Each is
727 labelled by its centre, so a feature at a given step is located to within half a window; the window is
728 not causal. The statistic reported per window is PR^{det} , the detrended effective rank of equation (3).
729 The total displacement is recorded beside it, since a rank can fall simply because the trajectory has
730 stopped moving.

731 G.2 The effective-rank collapse, run by run

732 Table 14 states the measurement of section 7.2 for every run. The plateau is the median over the 4 000
733 steps before the generalisation step; the dip is the minimum within $\pm 4 000$ steps of it; the column
734 headed “at” gives the position of that minimum relative to it. The two runs without weight decay
735 have no generalisation step, so each is measured over the window of the run it matches.

736 The last column is the largest value the statistic reaches after its dip. This is the sense of “re-expands”
737 in section 7.2. In parameter space that value exceeds the plateau by an order of magnitude in the
738 modular runs. In function space it returns to about the plateau and exceeds it in two of the four runs.

Table 14: Plateau, dip and recovery per run, in function space and in parameter space. Depth is plateau divided by dip. The two rows without weight decay in each block are measured in the aligned window; their deepest fall anywhere in their own run is 3.15 and 7.06 in function space and 1.66 and 2.37 in parameter space, the comparison section 7.2 also reports.

space	run	plateau	dip	depth	at	post-dip max
function	mod_wd1	24.24	3.80	6.38	+135	23.99
	mod_wd1_s43	33.02	3.86	8.55	455	27.80
	mod_wd1_s44	27.16	3.40	7.99	+515	30.82
	s5_wd1	21.04	2.62	8.03	1413	23.53
	mod_wd0	5.50	3.43	1.60	—	—
	s5_wd0	11.87	9.09	1.31	—	—
parameter	mod_wd1	3.01	1.09	2.76	+1 035	36.12
	mod_wd1_s43	2.57	1.15	2.24	1155	37.34
	mod_wd1_s44	2.71	1.13	2.39	+1 315	37.44
	s5_wd1	7.14	1.39	5.14	1513	12.46
	mod_wd0	29.93	27.70	1.08	—	—
	s5_wd0	5.57	4.49	1.24	—	—

H The resolution of a windowed statistic

H.1 The full-batch setting

Repeating the direct measurement on the perceptron setting gives a result that must be read as a limit of the measurement rather than as a negative finding. Under full-batch gradient descent every window of every run has a participation ratio between 1.000 and 1.362 on the statistics the comparison uses.

A deterministic update makes the trajectory a smooth curve. Over a 600-step window such a curve is straight to within the measurement precision, so the statistic reports the curvature of the trajectory at the scale the window sets, where there is almost none. The statistic has no range in which a collapse could appear. Its absence is therefore not evidence that no collapse occurred.

Repeating the measurement at a range of window lengths settles the question. Two runs of 150 000 steps, a generalising one and its label-matched partner, are measured over windows spanning 600 to 118 000 optimiser steps. Each window holds the same sixty samples, spread further apart as the window lengthens rather than supplemented by new ones, so that the bound on the statistic and its noise floor stay exactly where the published measurement put them. Table 15 gives the result. The statistic does leave its floor, at a scale of about 10^4 steps, more than an order of magnitude above the published window. The absence of a feature at the shorter window was therefore a limit of resolution rather than a fact about the run.

The feature that appears at the longer window is not a signal that was previously concealed. One localised maximum appears, of the same size in both runs (table 15). Function space agrees, with more range and no more separation. Nor can that maximum be placed relative to memorisation or to generalisation: a window this long cannot place a feature to better than the interval between memorisation and generalisation.

At the window where the statistic has range, it does not separate the generalising run from its label-matched control. Beyond 30 000 steps the ordering even reverses and the non-generalising run moves further, because the window is then a large fraction of a run whose weight norm is still climbing, so that the statistic tracks that drift. The full-batch trajectory therefore carries no signature of generalisation in its effective rank at any window length between 600 and 118 000 steps. The collapse of section 7.2 belongs to the regularised stochastic setting.

Table 15: Whether the full-batch statistic has any range at longer windows. Detrended position participation ratio in parameter space against window length, full-batch gradient descent at $p = 97$, over all placements of the window, sixty samples per window at every length.

window (steps)	generalises		control, never generalises	
	median	max	median	max
600	1.002	1.035	1.001	1.011
2 950	1.066	1.258	1.003	1.087
5 900	1.027	1.534	1.005	1.360
11 800	1.010	1.951	1.020	1.826
29 500	1.006	2.047	1.112	1.735
59 000	1.025	1.551	1.345	1.912
118 000	1.231	1.626	2.009	2.271

Introducing mini-batches confirms this. The same statistic then ranges over more than a decade while the generalisation step moves by a single logged step. Any claim based on a trajectory participation ratio must therefore be controlled for batch size and noise scale before it is attributed to learning.

H.2 Declines of the estimate that are not generalisation

A decline of the estimate is not specific to generalisation. In a long run at $p = 211$ with no weight decay, which memorises early and ends at chance validation accuracy, a decline is detected tens of thousands of steps after memorisation. The ramped-gain experiment of section 6.3 supplies a mechanism for it that needs no dimension at all: a ramped observer gain moves the estimate by about a component, while the parameter norm of this run roughly triples and acts on the estimate through the observer.

The trajectory of that run does not simplify. Its diagnostics place it in the transient regime (figure 7), where the estimate is not a dimension. The direct measurement of section 7.2 is an effective rank of the trajectory itself; it declines in four generalising runs and in neither of the two without weight decay. The estimate on a log, at the frozen window and in these settings, is neither a dimension nor an effective rank; it declines in at least one run that does not generalise.

Section 7.3 does not make the decline specific. The estimate on a log does decline where the trajectory does, in four runs and in neither of the other two, but at a window two orders of magnitude shorter than the one used here, under a configuration chosen for that window. Nor is a run's failure to generalise within the allotted steps evidence that a decline in it was a false alarm.

H.3 Resolution of a windowed estimate on a training log

A change within one run might still be read where the level may not be. That possibility is testable on the seven extended runs. The frozen configuration takes a window of a third of each record, 39 990 optimiser steps, advanced by 10 000 steps at a time. That gives nine windows per run, localising a feature only to $\pm 19\,995$ steps.

The runs therefore cannot be aligned on their generalisation steps as figure 2 aligns them. No two runs share an offset, so no mean over runs can be formed.

The one run with a window centre near its generalisation step falls by 3.54 components over the following stride, more than twice the 1.53-component floor that a ramped gain alone produces (section 6.3). The windows in which the estimate falls are the windows in which the parameter norm is itself moving fastest, so the fall is the very ramped gain the floor is built from. The other two runs show nothing.

Nor is any of this specific to generalisation. Across the fifty-six stride-to-stride changes (figure 9), six of twenty-four exceed that floor in the runs that generalise and six of thirty-two in the runs that do not. The largest change in the set, 6.36 components, is in a run that does generalise, but it occurs nearly ninety thousand steps after its generalisation step.

As in appendix H.1, the statistic has no range in which a feature confined to the transition could appear. Testing the claim requires a window set in units of the width of the transition, which for these runs is hundreds of steps. That means re-running the estimator, not re-reading the existing records.

The matched-window re-run. Section 7.3 performs that re-run on the six runs of figure 2 rather than the seven of this appendix: the question needs a paired comparison and only those six had their trajectory stored. Each window is 60 samples, centred on the midpoint of a window of section 7.2.

The depth of the fall is measured by D , the median estimate over the three thousand steps before the transition divided by the minimum over the window that brackets it. Its significance is assessed against the surrogates of section 7.3 [Theiler et al., 1992, Schreiber and Schmitz, 1996]. Table 16 gives both.

The frozen configuration does not fit. At 60 samples the frozen delay span of 76 exceeds the window, so some other configuration has to be used. The grid crosses $E_{\max} \in \{4, 6, 10\}$, $\tau \in \{1, 2, 4\}$, $m \in \{5, 20\}$ and both Theiler rules, giving thirty-six cells, of which twenty-six are long enough to return a value. Twenty-four of those separate the four generalising runs from the other two. The headline cell is named by the rule of section 7.3, which does not choose the best cell: a neighbouring cell separates the runs more sharply.

The diagnostics do not survive the shortening. The identifiability ratio needs an embedding of $2E_{\max}$, so at the headline cell the Theiler exclusion leaves too few usable points to compute one. Where the ratio can be computed at all, the two Theiler rules disagree about whether it passes; it settles nothing.

Linear detrending removes the effect. Repeating every cell on a linearly detrended window, the removal PR^{det} performs coordinate by coordinate, leaves no cell that separates the generalising runs from the rest. Detrending a sketch and detrending a scalar are not comparable: on a 1024-dimensional sketch a per-coordinate linear detrend removes one direction out of many, whereas on a scalar series it removes most of the series.

The observer matters. The result above is on the parameter norm. On the training loss, ten of the twenty-six cells separate them, though the surrogate comparison passes in the S_5 run only.

The two runs without weight decay have almost no dynamic range. At this window the baseline belongs to the configuration rather than to the run: on a standardised straight line equation (4) returns about 2.77, close to the pre-transition level of every run. The parameter-norm log of the modular run is straight to under a per cent of its own spread across the whole comparison interval, so its estimate spans three thousandths of a component, against a range some three orders of magnitude larger in the generalising runs. Its $D = 1.00$ therefore records that the observer offered the estimator no local structure to respond to. The S_5 run is a weaker case of the same thing.

The D and quantile columns of these two runs may not, then, be treated as informative dimensional negatives. The four generalising runs are untouched by this. So is the surrogate comparison, in which each run is measured only against surrogates of itself.

Table 16: The fall of the log estimate at generalisation, measured against shape-preserving surrogates. The matched-window re-run of section 7.3 on the parameter-norm log, at the headline cell $E_{\max} = 10$, $\tau = 1$, $m = 20$, embedding-span Theiler. D is the median estimate before the transition divided by the minimum across it; the quantile is the rank of D among the same statistic at every window centre after memorisation; the floor is where that minimum falls, relative to t_{gen} . The p columns are the fraction of surrogates, at three smoothing lengths, whose D is at least the observed one.

run	D	quantile	floor	p_{101}	p_{201}	p_{401}
modular, seed 42	2.96	0.99	−65	0.025	0.025	0.025
modular, seed 43	1.98	0.93	255	0.050	0.050	0.025
modular, seed 44	2.14	0.92	+315	0.050	0.075	0.025
S_5 composition	2.12	0.89	1313	0.050	0.050	0.050
modular, no decay	1.00	0.46	−865	0.950	0.950	0.925
S_5 , no decay	1.04	0.38	1613	0.975	0.975	1.000

838 The linear participation ratio and the roughness were put through the same comparison. The linear
839 participation ratio responds less than the estimator, the roughness by more than an order of magnitude.




Figure 9: How coarsely a frozen-window estimate localises a feature. The windowed estimate on the parameter norm for the seven extended runs, under the frozen configuration of section 7.1, plotted at the window centres. Colour is the regime of section 3.3. (a) Every window, with the generalisation steps marked on the axis. (b) The generalising runs aligned on their own generalisation step. (c) Every stride-to-stride change, one sub-row per run. The grey band in (b) and (c) is the floor a ramped gain alone produces (section 6.3).

840 I The Theiler exclusion

841 I.1 The exclusion cap at the twenty-component configuration

842 The rule this study asks for is $W_T = \max((E - 1)\tau, t_{\text{acf}})$: at least the full span of a delay vector,
843 so that two neighbours cannot be near each other merely because their windows overlap. The
844 implementation caps W_T at 150 samples.

845 At the eight-component configuration the embedding span is 76, so the cap binds only where
846 the autocorrelation time exceeds 150. That never happens on the constructed systems, whose
847 autocorrelation time is a few samples, but it happens on five of the seven parameter-norm records
848 of section 7. Those five were scored at $W_T = 150$ rather than at the value the rule asks for, in the
849 direction that inflates the estimate. At the twenty-component configuration the span is 624 and the
850 cap binds unconditionally. It is why section 5.2 marks that result exploratory.

851 Table 17 measures the cost of that departure, by re-scoring the stored trajectories at both exclusions.

Table 17: The cost of the exclusion cap. The twenty-component system re-scored at the capped exclusion and at the one the configuration asks for. Seven ranks, three seeds, the three observers with the lowest error, median over seeds; the error is against the measured PR^{pos} and is not comparable with table 4, which covers more observers.

r	measured PR^{pos}	MG at $W_T = 150$	MG at $W_T = 624$
2	2.00	1.74	1.81
5	5.00	4.16	5.15
8	8.00	6.28	6.71
11	11.00	9.54	10.26
14	14.00	9.81	9.37
17	17.00	14.40	14.19
20	19.99	15.92	16.60
mean absolute error		2.16	1.89
Spearman ρ		1.00	0.96

852 The ordering of the ranks, all that section 5.2 claims at this range, is nearly unaffected, a single
853 inversion apart. The mean absolute error falls when the protocol is enforced, so the cap was not
854 making the reported result look better than it is. The saturation above about six components is present
855 at both settings and is therefore not an artefact of the exclusion.

856 The same cap has a second effect, which the identifiability ratio inherits. The numerator of the ratio
857 is computed at $2E_{\max}$ with τ held fixed, so its delay span exceeds the cap while the denominator's
858 does not.

859 Five of the seven parameter-norm records of section 7 are unaffected by this. The asymmetric cases
860 are the other two and, more consequentially, the constructed systems of section 5; the reference band
861 of section 6.4 was therefore itself measured with 76 against 150. A ratio built from two differently
862 excluded estimates is not the quantity section 3.4 describes.

863 I.2 The Theiler exclusion and the dimension of a transient

864 Section 3.3 claims that recurrence is a requirement of this protocol and not of the estimator. The claim
865 is testable on the excitation cases of section 5.2 by varying the Theiler exclusion W_T alone, from
866 zero up to the value the frozen rule gives, with the estimator otherwise at the frozen configuration
867 (table 18).




Figure 10: The three regimes as raw scalar logs. Each panel is the parameter norm, standardised to zero mean and unit variance as in Algorithm 1. Left, the constrained head of section 5.2 under a fast quasiperiodic drive, whose path lies on a torus. Middle, the same system under plain gradient descent, which decays over its whole record. Right, a transformer grokking modular addition.

Table 18: The estimate against the Theiler exclusion, median over observers, seeds and windows. The transient's estimand is 1: at $W_T = 0$ the estimator returns it, whereas the exclusion removes it. The fast recurrent case is unchanged under the same variation. The last column removes the cap of 150 and applies the autocorrelation rule, some 1 600 samples here.

W_T	0	2	5	20	100	150	uncapped
recurrent, fast	4.83	4.82	4.82	4.82	4.80	4.80	4.80
recurrent, slow	2.06	3.02	3.22	3.27	3.27	3.28	3.27
transient	1.20	1.71	2.38	5.51	20.50	29.22	173.19

868 *The value at $W_T = 0$ is the right one.* On a uniformly sampled curve the nearest neighbours of
869 a point are its immediate temporal neighbours, for which equation (4) returns 1.199; the transient
870 returns that value in every cell. The excess over unity is the lattice bias of a deterministically sampled
871 curve.

872 *The exclusion removes the near neighbours, not the neighbours.* The estimator fills its twenty slots at
873 every exclusion; only the distances change. Before any exclusion the nearest neighbours of a point lie
874 thousands of samples away under the fast quasiperiodic drive, but six samples away on the transient
875 (figure 10). Excluding 150 samples therefore costs the fast drive almost none of its close pairs and
876 the transient all of them (figure 11). The distances the transient retains are then nearly equal, their

log-ratio sum collapses, and the reciprocal in equation (4) diverges. An exponential decay with no optimiser and no data behind it reproduces both ends.

The asymmetry is general; the value 29 is not. Recurrence is not always needed even here. Under the fast drive $W_T = 0$ already returns r , because that case is not oversampled; it is the slow drive that the exclusion repairs, its nearest neighbours lying some six hundred samples away. A transient also need not have dimension one: a decaying *oscillation* correctly returns 2.36 at every exclusion.




Figure 11: Why the exclusion sets the value on a transient and leaves a torus alone. The excitations of table 18, median over ranks, seeds and windows. In (b) a neighbour counts as a return when it is no further away than the nearest neighbours an unexcluded search finds, so a fall to zero leaves the estimator nothing but distant points.

A recurrence-based diagnostic. The ratio $\hat{d}_{MG}(W_T=150)/\hat{d}_{MG}(W_T=20)$ is unity under both recurrent drives and above five on the transient. Unlike the trend-crossing count, this ratio tests recurrence in the reconstructed space rather than non-monotonicity of the series. That is the gap section 3.4 leaves open, which appendix K shows to matter on real runs. The principled alternative is older than either: a space-time separation plot [Provenzale et al., 1992] shows directly how neighbour distance depends on separation in time. Choosing the exclusion by such a measurement rather than by a rule is the obvious improvement to the protocol.

J The training runs

Table 19 lists every training run the paper draws on. Two architectures appear. **T** is the one-layer transformer of Nanda et al. [2023]: width 128, four heads, no layer normalisation, about 227 000 parameters, trained in double precision with AdamW. **P** is the two-layer quadratic perceptron of Gromov [2023]: width 500, 145 500 parameters at $p = 97$, trained by full-batch gradient descent with no regularisation.

Nanda et al. [2023] train at full batch, whereas we use mini-batches, which makes the transformer setting stochastic rather than deterministic. In the perceptron setting we take the tasks and the representability criterion of Doshi et al. [2024] but not their optimiser, running plain full-batch gradient descent on mean squared error instead, which makes those runs transient. Separately, the S_5 task [Power et al., 2022] and its weight decay lie outside the configuration of Nanda et al. [2023]. The two S_5 runs also step the optimiser twice per batch, so their effective learning rate and decay are about double the nominal ones.

Definition of memorisation and generalisation. The memorisation step is the first logged step at which the training accuracy reaches 95 %; the generalisation step is the first at which the validation accuracy does. Neither requires the threshold to be held thereafter. The extended reruns of section 7.1 use the stricter rule instead, requiring the threshold to hold over a sustained block.

Table 19: Training runs used in this paper. α is the training fraction, “final” the validation accuracy at the end of training. A dash in the generalisation column marks a run that had not generalised when training ended; such runs are censored observations, not established negatives.

run	task	arch	wd	α	steps	mem.	gen.	final	used in
<i>Trajectory sketch attached</i>									
mod_wd1	mod. add. $p=113$	T	1.0	0.30	20k	1 470	13660	1.00	section 7.2
mod_wd1_s43	mod. add. $p=113$	T	1.0	0.30	20k	1 580	6640	1.00	section 7.2
mod_wd1_s44	mod. add. $p=113$	T	1.0	0.30	20k	1 410	3 680	1.00	section 7.2
s5_wd1	S_5 composition	T	0.2	0.50	15k	1 280	6685	1.00	section 7.2
mod_wd0	mod. add. $p=113$	T	0.0	0.30	20k	630	—	0.03	section 7.2
s5_wd0	S_5 composition	T	0.0	0.50	15k	1 015	—	0.01	section 7.2
<i>Extended reruns, 120 000 steps</i>									
gropos_s0	mod. add. $p=113$	T	1.0	0.30	120k	1 450	5070	1.00	section 7.1
lowdata15_s0	mod. add. $p=113$	T	1.0	0.15	120k	620	—	0.58	section 7.1
lowdata15_s1	mod. add. $p=113$	T	1.0	0.15	120k	630	106 380	1.00	section 7.1
lowdata15_s2	mod. add. $p=113$	T	1.0	0.15	120k	660	—	0.01	section 7.1
lowdata20_s0	mod. add. $p=113$	T	1.0	0.20	120k	800	38380	1.00	section 7.1
wd0_s0	mod. add. $p=113$	T	0.0	0.30	120k	610	—	0.09	section 7.1
wd0_s1	mod. add. $p=113$	T	0.0	0.30	120k	630	—	0.07	section 7.1
p211_wd0	mod. add. $p=211$	T	0.0	0.16	200k	1 800	—	0.04	section 7.1
<i>Perceptron, $p = 97$, full batch, no regularisation</i>									
a_add	$n + m$	P	0	0.5	100k	9230	11760	1.00	section 7.1
a_mul	$n \cdot m$	P	0	0.5	100k	9100	11560	1.00	section 7.1
a_sub	$n - m$	P	0	0.5	100k	9110	11540	1.00	recorded only
a_sq_sum	$n^2 + m^2$	P	0	0.5	100k	8440	10470	1.00	recorded only
a_sum_sq	$(n + m)^2$	P	0	0.5	46k	7350	8190	1.00	recorded only
x_mix_quad	mixed quadratic	P	0	0.5	100k	9600	—	0.01	figure 7
x_no_grok	$n^3 + nm^2 + m$	P	0	0.5	100k	9860	—	0.01	section 7.1
g_p1	$(4n_1 + n_2^3)^3$	P	0	0.5	100k	7470	12290	1.00	section 7.1
g_p1x	$+ n_1 n_2$	P	0	0.5	100k	10140	—	0.01	section 7.1
g_p2	$(2n_1 + 3n_2)^4$	P	0	0.5	100k	6470	7680	1.00	section 7.1
g_p2x	$- n_1^2$	P	0	0.5	100k	9730	—	0.02	section 7.1
g_p3	$(5n_1^3 + 2n_2^4)^2$	P	0	0.5	100k	5970	6600	1.00	section 7.1
g_p3x	$- n_2$	P	0	0.5	100k	9420	—	0.74	section 7.1

Four of the perceptron runs are repeated with the trajectory sketch attached, for the measurement of appendix H.1: a_add and x_no_grok from the table above, together with the label-matched pair g_p1 and g_p1x. Each is run at full batch and again with mini-batches of 512, two of them once more at $p = 23$. The eight runs at $p = 97$ reproduce the memorisation and generalisation steps of the rows above, so appendix H.1 compares two noise levels on matched learning.

One row needs a note. g_p3x ends far above the majority-class baseline of 1.3% while being provably outside the class of functions the architecture can represent.

K Descent at the edge of stability

Descent at the edge of stability [Cohen et al., 2021] is gradient descent at a fixed rate: the top Hessian eigenvalue sharpens until it reaches $2/\eta$ and then stays there, while the loss descends over long stretches and not over short ones. Such a run is deterministic and is not a plain transient, so an undriven trajectory could satisfy the conditions of section 3.3.

Setting. We run full-batch descent without weight decay on the quadratic perceptron of appendix J, over a range of learning rates and two seeds (table 20). The log is written at every optimiser step, because the oscillation is the two-cycle of the unstable mode and its period is a few steps. The sharpness is measured by power iteration on Hessian-vector products, taking the largest algebraic eigenvalue, because the Hessian is indefinite early in training and plain power iteration would then return a negative sharpness.

Table 20: Oscillation and recurrence in undriven descent at the edge of stability. Full-batch descent by learning rate, both seeds. The stability ratio is $\eta\lambda_{\max}/2$ over the second half of the run; *rises* is the fraction of post-transition steps on which the loss increases; *return* is the nearest neighbour surviving a 150-sample exclusion divided by the distance the trajectory travels during that exclusion, for which a two-frequency torus scores 0.008.

η	stability ratio	rises	return	outcome
1×10^5	0.239, 0.238	0.000	1.007	monotone
3×10^5	0.758, 0.745	0.000	1.007, 1.007	monotone
1×10^6	0.969, 0.971	0.495, 0.490	0.104, 0.079	at the edge
1.5×10^6	0.974, 0.970	0.500, 0.500	1.012, 1.011	at the edge
2×10^6	0.977, 0.978	0.500, 0.500	0.120, 1.013	at the edge
2.5×10^6	—, 0.988	—, 0.500	—, 0.034	one seed diverges
2.8×10^6	0.992, —	0.503, —	0.022, —	one seed diverges
3×10^6	—	—	—	diverges

Oscillation is not recurrence. Eight runs reach the edge and all eight oscillate, the loss rising on one step in two, but only five of them return close to states already passed through. The return column of table 20 separates those two groups by a factor of ten with nothing in between.

Performance of the diagnostics. The band of section 3.4 rejects all eight runs at the edge, on the training loss and on the parameter norm alike, whether or not they recur. Nor does the estimate separate the runs that return from those that do not: their ranges on the training loss overlap. The trend-crossing count tests non-monotonicity and not recurrence, so it passes every run at the edge, while ρ_{ident} stays far outside the band in both groups.

Dependence on the logging stride. Written at the stride every other full-batch log here uses, most of the phenomenon disappears (figure 12): over these eight runs the median trend-crossing count falls from 71 to 6. Whether a log looks like a transient is thus in part a property of how often it was written.




Figure 12: Whether full-batch descent reaches the edge of stability. (a) The stability ratio along training, one seed per rate. (b) Sixty consecutive steps at one rate at the edge, against the same series at a stride of ten: the two-cycle is removed by that stride rather than blurred by it. (c) The fraction of steps on which the loss rises, against the logging stride.

L The ceiling: embedding dimension against record length

Section 5 reports that the estimator saturates above about eight components. Two explanations were available: the embedding condition $E > 2d$ and the finite-record bound $2 \log_{10} N$ [Eckmann and Ruelle, 1992]. Each is insensitive to the variable the other names, so we sweep $E_{\max}$ at a fixed record length and the record length at fixed $E_{\max}$, over $r \in \{2 \dots 20\}$ and three seeds with everything else frozen (table 21). Each record length is a *prefix* of one longer record rather than a resampling of a fixed span, leaving the delay lag its meaning in periods. Every trajectory is bit-identical to the published one over their common prefix.

Table 21: Whether the ceiling follows the embedding condition or the finite-record bound. Two sweeps, three seeds per cell. *Tracking* is the largest r at which the median estimate stays within one component of the truth; *level* is the median estimate at $r = 20$; *slope* is $d\hat{d}_{MG}/dr$ over $r \geq 8$, which is 1 if the estimator still tracks the rank and 0 if it has stopped.

$E_{\max}$ at $N = 8000$	10	14	20	28	40	56	
tracking	6.5	7.3	8.1	10.3	10.9	11.2	
$E_{\max}/2$ predicts	5	7	10	14	20	28	
level at $r = 20$	6.9	8.1	9.8	11.3	13.2	14.9	
slope	0.08	0.12	0.20	0.29	0.41	0.54	
N at $E_{\max} = 20$	10^3	$2 \cdot 10^3$	$4 \cdot 10^3$	$8 \cdot 10^3$	$1.6 \cdot 10^4$	$3.2 \cdot 10^4$	$6.4 \cdot 10^4$
tracking	6.7	8.3	8.5	8.1	8.9	9.5	9.5
$2 \log_{10} N$ predicts	6.0	6.6	7.2	7.8	8.4	9.0	9.6
level at $r = 20$	9.0	9.3	9.4	9.8	9.9	10.2	10.5
slope	0.22	0.17	0.17	0.20	0.19	0.19	0.22

Outcome. The level at high rank is set by the embedding: a sixty-fourfold increase in the record adds 1.5 components, whereas a sixfold increase in $E_{\max}$ adds eight. The tracking ceiling follows neither variable. It never passes 11.2, yet at a quarter of the record the rest of the study uses it still tracks to 8.3, where $2 \log_{10} N$ allows only 6.6. A bound set by the record length cannot produce that, so [Eckmann and Ruelle \[1992\]](#) is not the explanation here. Nor is $E > 2d$ in the form we used it: the ceiling rises by a tenth of a component per unit of $E_{\max}$ rather than by a half: $E_{\max}/2$ happens to be right near $E_{\max} = 14$ and is wrong at both ends.

Fitting the level at high rank separates the candidates (table 22): neither hypothesis fits, nor does their pointwise minimum, whereas a form logarithmic in both variables fits an order of magnitude more closely than the spread of the data itself.

Table 22: Whether either candidate explanation of the ceiling accounts for the measured level. Candidate forms fitted at $r \geq 16$ over every measured $(E_{\max}, N)$ cell; the spread of the fitted quantity itself is 2.00 components, so a form with a larger error than that describes nothing.

form	root-mean-square error
$E_{\max}/2$, the embedding condition	8.73
$2 \log_{10} N$, the finite-record bound	3.68
the pointwise minimum of the two	3.70
$8.61 \log_{10} E_{\max} + 0.90 \log_{10} N - 5.31$	0.19

A purely linear statistic reproduces the whole dependence. Across the $E_{\max}$ sweep PR_{delay} rises over the same range as MG; across the record sweep it is flat; and a single logarithm in $E_{\max}$ fits it more tightly than any form fits anything else in this appendix. Most of the ceiling is therefore the number of components the delay window resolves *linearly*.

The unexplained tracking limit. We cannot say why the tracking ceiling lies near eleven components. It does not follow the measured effective rank, which equals r throughout, nor does any form we fitted account for it.

The identifiability ratio of section 3.4 does register it. At $r = 20$ the ratio stays above the band at every record length, so a longer record never makes an unidentifiable estimate identifiable; it falls into the band as $E_{\max}$ grows, at about the rank where tracking stops.

M Two controls separating the estimator from roughness

Roughness is $\text{std}(\Delta x)/\text{std}(x)$. The controls below separate it from the estimate in both directions, each given its pass criterion before it was run.

968 *The number of components changes and the roughness does not.* Both systems in this control are a
969 sum of four sinusoids. In the first the four frequencies are rationally independent, so the path lies
970 on a four-dimensional torus and $d_{\text{act}} = 4$. In the second they are harmonics of one irrational base
971 frequency: a single phase fixes all four sinusoids at once, the path is a circle, and $d_{\text{act}} = 1$. The base
972 frequency of the second system is then set by bisection so that the two have the *same* roughness, to a
973 part in 10^{15} . A scheduled run passes between the two through ramped envelopes, leaving unscored
974 every window a ramp reaches into (figure 13).




Figure 13: A change in the number of independent phases, tracked as it happens. The harmonic control of this appendix, scheduled from four phases to one and back; median over eight seeds, band and their range. Dashed in (a), the number of phases, drawn only where no ramp reaches into the window. The linear statistic of (c) moves in the opposite direction.

975 *The roughness changes and the number of components does not.* Five monotone observers $g(x) =$
976 $x + ax^p$ are applied to one four-phase system. Each is invertible, so it may bend the recorded set but
977 cannot add or remove a coordinate. Each also changes the local scale $g'(x)$ and with it the spacing of
978 consecutive samples.

Table 23: Both controls, median over eight seeds. Above the rule, two systems matched in roughness and differing in the number of independent phases. Below it, one four-phase system under five invertible maps. PR_{delay} is the linear participation ratio of the delay covariance, included because it follows the roughness rather than the active dimension in both halves.

	d_{act}	estimate	roughness	PR_{delay}
four independent phases	4	3.92	0.091	2.20
one phase, four harmonics	1	1.11	0.091	4.10
identity	4	3.92	0.091	2.20
$x + 0.25x^3$	4	3.96	0.098	2.32
$x + x^3$	4	3.96	0.110	2.58
$x + x^5$	4	4.01	0.146	3.61
$x + 0.1x^7$	4	4.02	0.166	4.28

979 Both criteria were fixed before the runs: that the estimate should follow the number of phases while
980 the roughness holds still, and that it should stay within a tenth of a component of four while the
981 roughness moves. Table 23 shows both hold.

982 **The limits of this control.** Both controls separate the statistics on a constructed system whose
983 phases are known. They say nothing about a training log, where the comparison against surrogates in
984 section 7.3 carries the claim instead.

985 Two findings count against the estimator. The linear participation ratio reads higher on the circle than
986 on the four-torus (table 23), so anyone treating it as a count of phases would order the two backwards.
987 On the circle the estimate is also not geometric: every neighbour surviving the exclusion there is a
988 sample from another pass around the circle. The local speed of the curve multiplies all of a point's
989 neighbour distances alike; speed cancels from the log-ratios of equation (4). The residue depends on
990 the arithmetic of the base frequency: two of the eight seeds drew a frequency whose multiples cluster

991 into near-duplicates; their estimate is reproducible from the fractional parts of those multiples alone,
992 with no signal and no neighbour search.

993 **N Compute cost and code availability**

994 **Compute cost.** Everything here runs in a few hours on one consumer GPU and one desktop CPU.
995 The largest GPU item is the edge-of-stability experiment of appendix K, which takes half an hour on
996 a single T4; the analysis is the more expensive half and runs entirely on CPU, a few core-hours per
997 experiment.

998 **Cost of the trajectory sketch.** At each logged step the sketch keeps 4096 floats against the full
999 parameter vectors of appendix J: a run whose raw trajectory would occupy 88 megabytes occupies
1000 2.5. We train the same configuration with and without the sketch attached, alternating the order so
1001 that warm-up cannot favour either. The measured difference is negative, which is a bound rather than
1002 a result: at the logging stride the transformer runs use, the cost of the sketch lies below the run-to-run
1003 variation of the machine. `sketch_cost.py` records the measurement.

1004 **Code availability.** All code, configurations and result files behind every number and figure in this
1005 paper will be released publicly on acceptance. The figures regenerate from the committed result files
1006 with a single command.